

CS162

Operating Systems and Systems Programming

Lecture 3

Processes (con't), Fork, System Calls

January 28th, 2020
 Prof. John Kubitowicz
<http://cs162.eecs.Berkeley.edu>

Recall: Four Fundamental OS Concepts

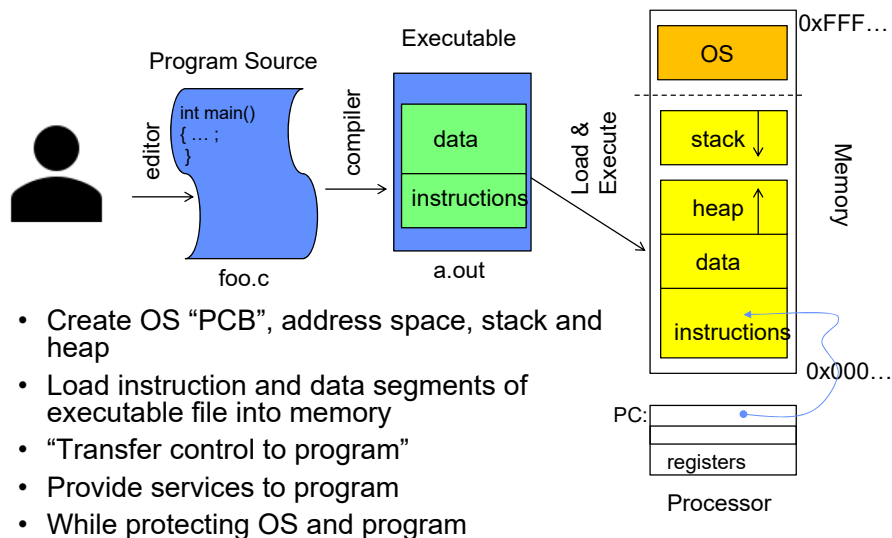
- **Thread: Execution Context**
 - Fully describes program state
 - Program Counter, Registers, Execution Flags, Stack
- **Address space (with or w/o translation)**
 - Set of memory addresses accessible to program (for read or write)
 - May be distinct from memory space of the physical machine (in which case programs operate in a virtual address space)
- **Process: an instance of a running program**
 - Protected Address Space + One or more Threads
- **Dual mode operation / Protection**
 - Only the “system” has the ability to access certain resources
 - Combined with translation, isolates programs from each other and the OS from programs

1/28/20

Kubitowicz CS162 ©UCB Spring 2020

Lec 3.2

Recall: OS Bottom Line: Run Programs



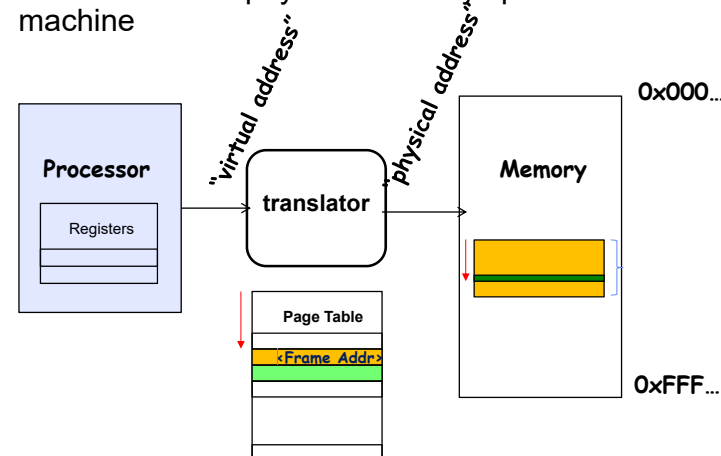
1/28/20

Kubitowicz CS162 ©UCB Spring 2020

Lec 3.3

Recall: Protected Address Space

- Program operates in an address space that is distinct from the physical memory space of the machine

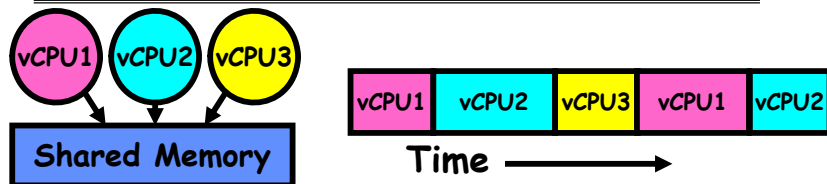


1/28/20

Kubitowicz CS162 ©UCB Spring 2020

Lec 3.4

Recall: give the illusion of multiple processors?



- Assume a single processor. How do we provide the illusion of multiple processors?
 - Multiplex in time!
 - Multiple "virtual CPUs"
- Each virtual "CPU" needs a structure to hold:
 - Program Counter (PC), Stack Pointer (SP)
 - Registers (Integer, Floating point, others...?)
- How switch from one virtual CPU to the next?
 - Save PC, SP, and registers in current state block
 - Load PC, SP, and registers from new state block
- What triggers switch?
 - Timer, voluntary yield, I/O, other things

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.5

Recall: The Process

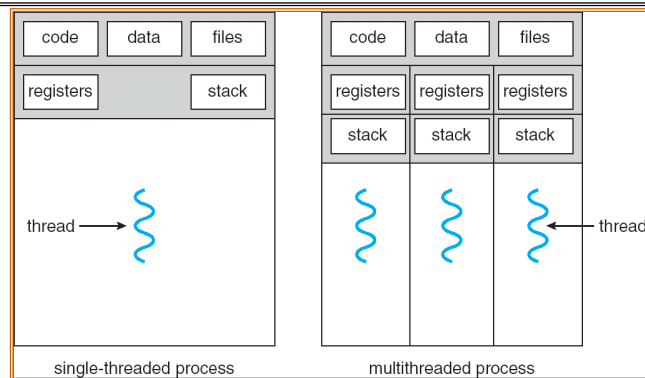
- Definition: **execution environment with restricted rights**
 - Address Space with One or More Threads
 - » *Page table per process!*
 - Owns memory (mapped pages)
 - Owns file descriptors, file system context, ...
 - Encapsulates one or more threads sharing process resources
- Application program executes as a process
 - Complex applications can fork/exec child processes [later]
- Why processes?
 - Protected from each other. OS Protected from them.
 - Execute concurrently [trade-offs with threads? later]
 - Basic unit OS deals with

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.6

Recall: Single and Multithreaded Processes



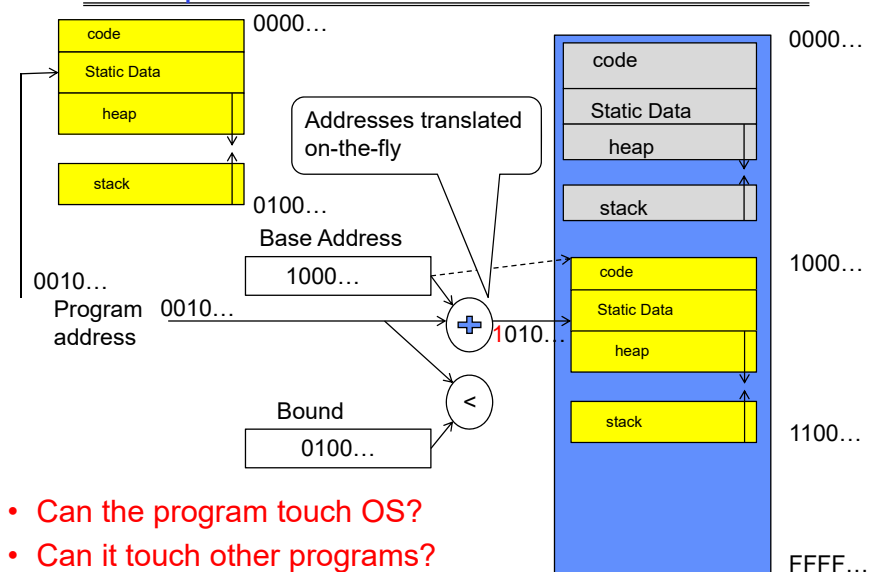
- Threads encapsulate concurrency: "Active" component
- Address spaces encapsulate protection: "Passive" part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.7

Recall: Simple address translation with Base and Bound



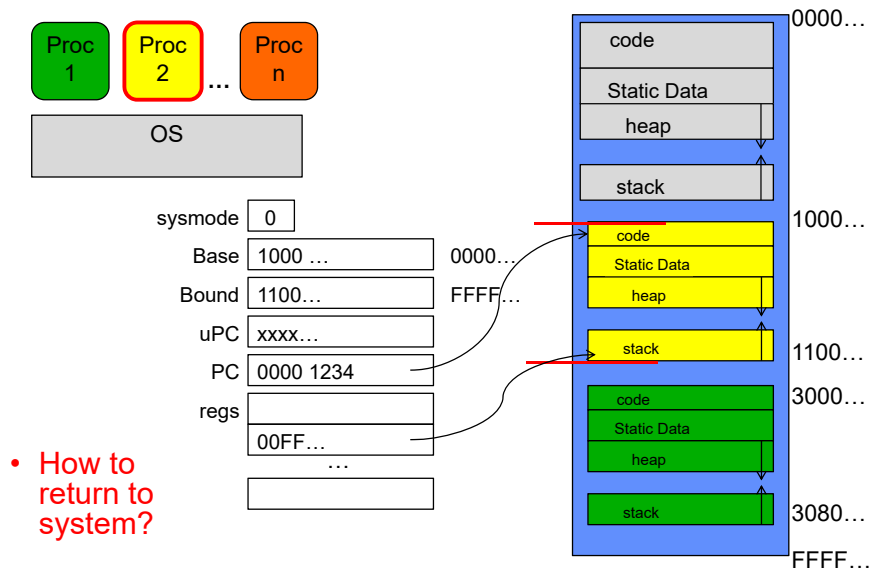
- Can the program touch OS?
- Can it touch other programs?

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.8

Simple B&B: User => Kernel

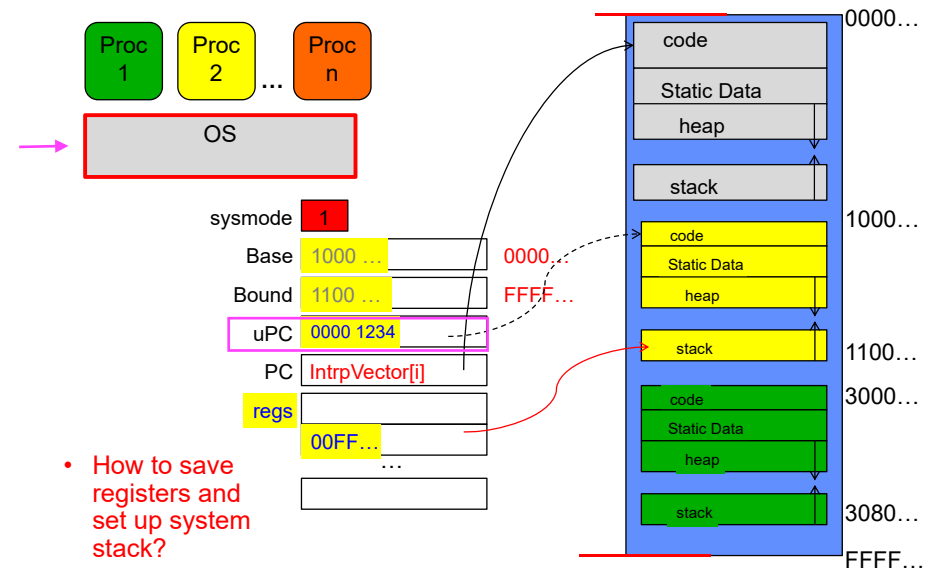


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.9

Simple B&B: Interrupt



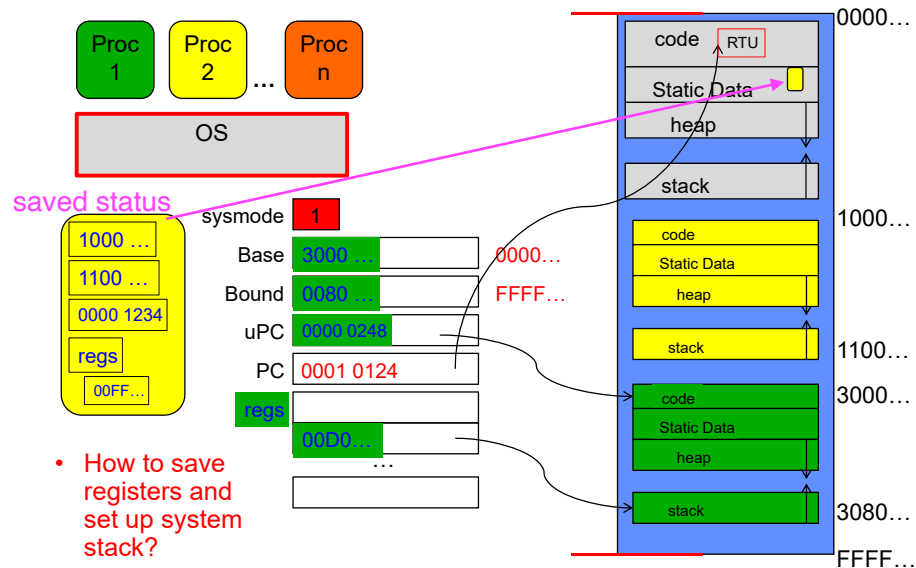
1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.10

IntrapVector decides where to go in the OS. Timer interrupt handler says ok yellow is done, let's do something else.

Simple B&B: Switch User Process

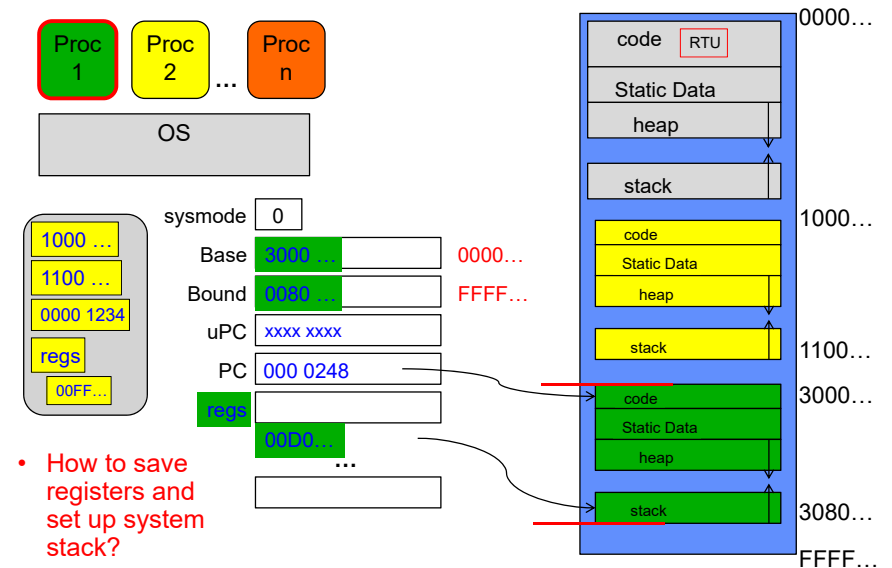


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.11

Simple B&B: "resume"



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.12

Is Branch and Bound a Good-Enough Protection Mechanism?

- **NO: Too simplistic for real systems**
- Inflexible/Wasteful:
 - Must dedicate physical memory for *potential* future use
 - (Think stack and heap!)
- Fragmentation:
 - Kernel has to somehow fit whole processes into contiguous block of memory
 - After a while, memory becomes fragmented!
- Sharing:
 - Very hard to share any data between Processes or between Process and Kernel
 - Need to communicate indirectly through the kernel...

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.13

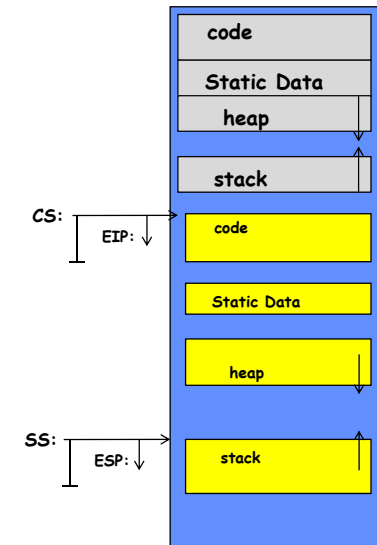
Better: x86 – segments and stacks

each segment has its own base and bound

Processor Registers

CS	EIP
SS	ESP
DS	EAX
ES	EBX
	ECX
	EDX
	ESI
	EDI

Start address, length and access rights associated with each segment

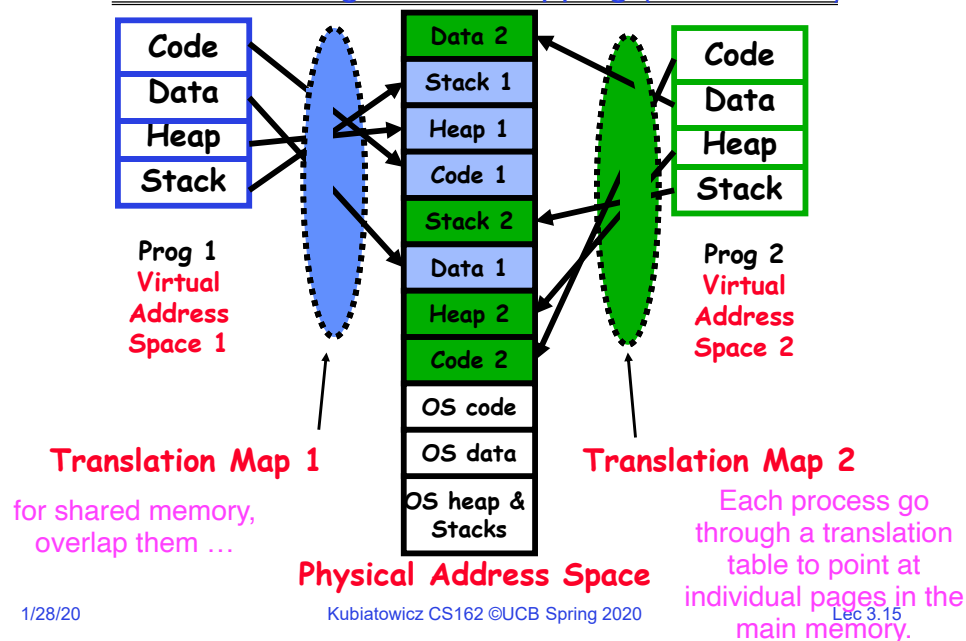


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.14

Alternative: Page Table Mapping (More soon!)

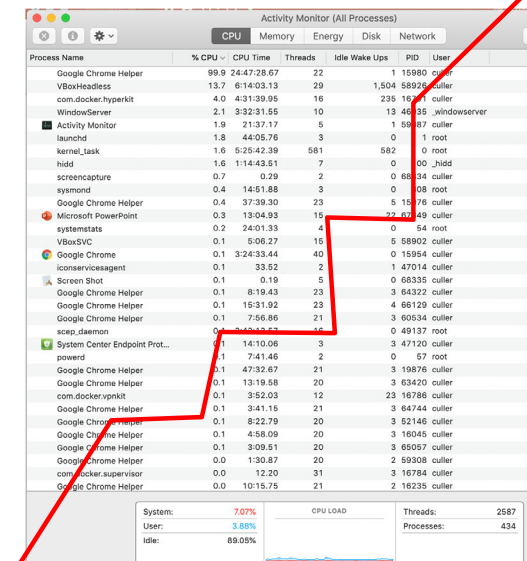


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.15

What's beneath the Illusion?



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

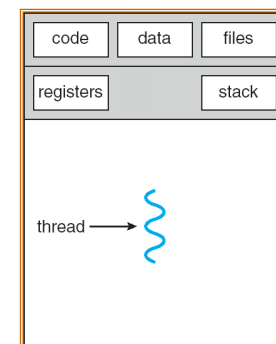
Lec 3.16

Today: How does the Operating System create the Process Abstraction?

- What data structures are used?
- What machine structures are employed?
 - Focus on x86, since will use in projects (and everywhere)

Starting Point: Single Threaded Process

- Process: OS abstraction of what is needed to run a single program
 1. Sequential program execution stream
 - » Sequential stream of execution (thread)
 - » State of CPU registers
 2. Protected resources
 - » Contents of Address Space
 - » I/O state (more on this later)

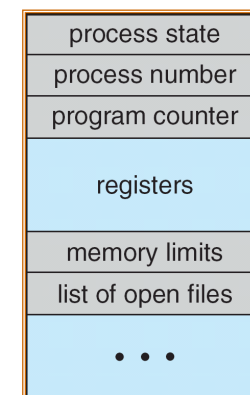


Running Many Programs

- We have the basic mechanism to
 - switch between user processes and the kernel,
 - the kernel can switch among user processes,
 - Protect OS from user processes and processes from each other
- Questions ???
 - How do we represent each process in the kernel?
 - How do we decide which user process to run?
 - How do we pack up the process and set it aside?
 - How do we get a stack and heap for the kernel?
 - Aren't we wasting a lot of memory?

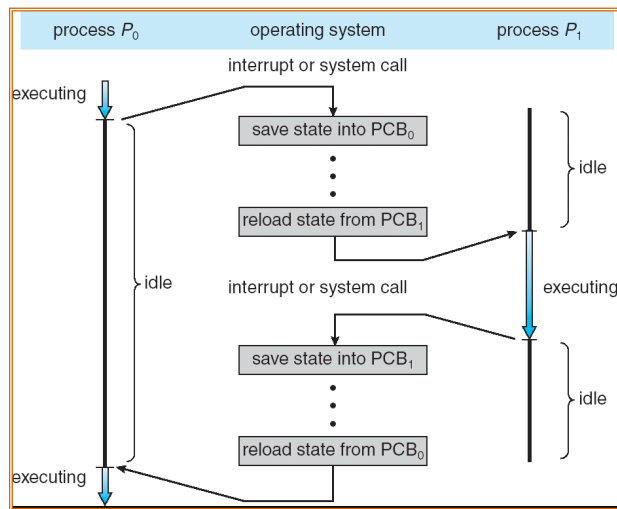
Multiplexing Processes: The Process Control Block

- Kernel represents each process as a process control block (PCB)
 - Status (running, ready, blocked, ...)
 - Register state (when not ready)
 - Process ID (PID), User, Executable, Priority, ...
 - Execution time, ...
 - Memory space, translation, ...
- Kernel *Scheduler* maintains a data structure containing the PCBs
 - Give out CPU to different processes
 - This is a Policy Decision
- Give out non-CPU resources
 - Memory/I/O
 - Another policy decision



Process
Control
Block

Context Switch

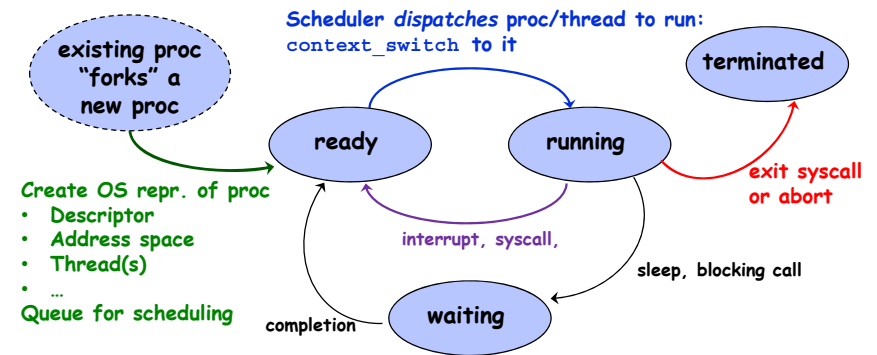


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.21

Lifecycle of a process / thread



- OS juggles many process/threads using kernel data structures
- Proc's may create other process (fork/exec)
 - All starts with init process at boot

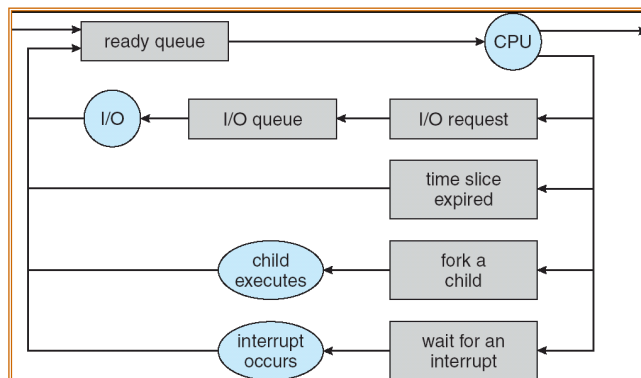
Pintos: process.c

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.22

Scheduling: All About Queues



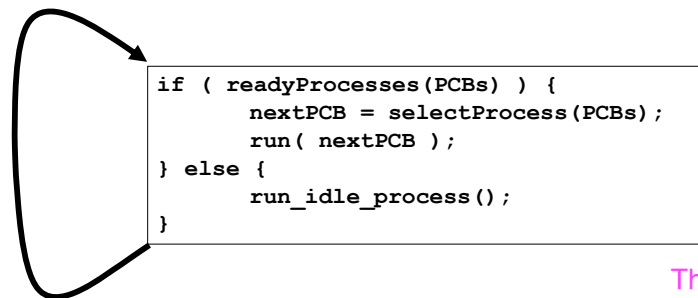
- PCBs move from queue to queue
- **Scheduling**: which order to remove from queue
 - Much more on this soon

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.23

Scheduler



- Scheduling: Mechanism for deciding which processes/threads receive the CPU
- Lots of different scheduling policies provide ...
 - Fairness or
 - Realtime guarantees or
 - Latency optimization or ..

The idle process typically tries to put the processor in a low-power status.

1/28/20

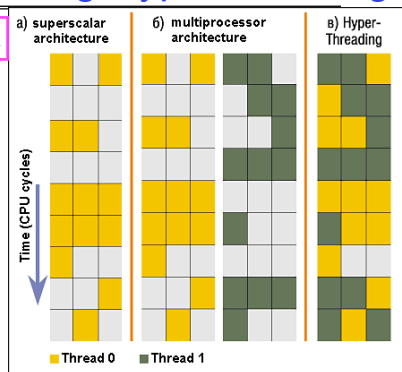
Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.24

Simultaneous MultiThreading/Hyperthreading

Hardware scheduling technique

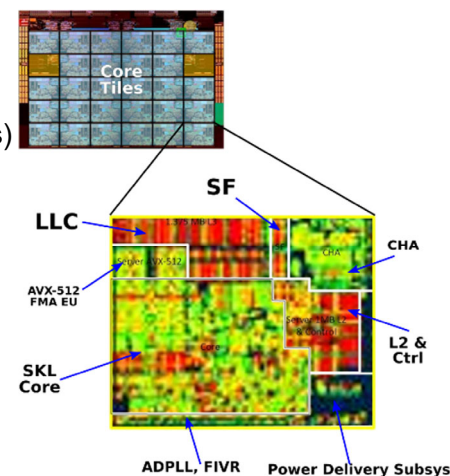
- Superscalar processors can execute multiple instructions that are independent.
- Hyperthreading duplicates register state to make a second "thread," allowing more instructions to run.
- Can schedule each thread as if were separate CPU
 - But, sub-linear speedup!
- Original technique called "Simultaneous Multithreading"
 - <http://www.cs.washington.edu/research/smt/index.html>
 - SPARC, Pentium 4/Xeon ("Hyperthreading"), Power 5



Colored blocks show instructions executed

Also Recall: The World Is Parallel

- Intel Skylake (2017)
 - 28 Cores
 - Each core has two hyperthreads!
 - So: 54 Program Counters(PCs)
- Scheduling here means:
 - Pick which core
 - Pick which thread
- Space of possible scheduling much more interesting
 - Can afford to dedicate certain cores to housekeeping tasks
 - Or, can devote cores to services (e.g. Filesystem)



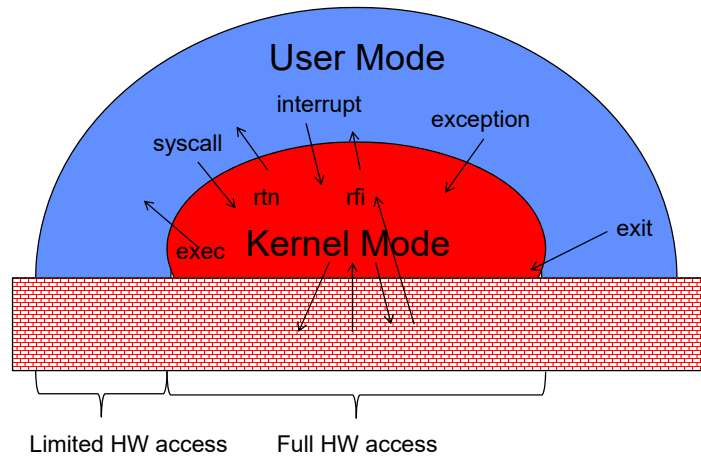
Administrivia: Getting started

- Kubiatowicz Office Hours:
 - 1-2pm, Monday & Thursday
- Homework 0 **Due Friday!**
 - Get familiar with the cs162 tools
 - configure your VM, submit via git
 - Practice finding out information:
 - » How to use GDB? How to understand output of unix tools?
 - » We don't assume that you already know everything!
 - » Learn to use "man" (command line), "help" (in gdb, etc), google
- **Should be going to sections now – Important information there**
 - Any section will do until groups assigned
- Class status: All regular students made it!
 - Concurrent enrollment will be added as possible. Over half will be admitted already. Perhaps more.
- **THIS Friday is Drop Deadline! HARD TO DROP LATER!**
 - If you know you are going to drop, please do so to leave room for others!

Administrivia (Con't)

- Group sign up via autograder form next week
 - Get finding groups of 4 people ASAP
 - Priority for same section; if cannot make this work, keep same TA
 - Remember: Your TA needs to see you in section!
- Midterm 1 conflicts
 - We will handle these conflicts after have final class roster
 - I know about one problem with Midterm 1 scheduling, and it can be dealt with. Have I missed any others?
 - Watch for queries by HeadTA to collect information

Recall: User/Kernel (Privileged) Mode



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.29

Three types of Kernel Mode Transfer

- Syscall
 - Process requests a system service, e.g., exit
 - Like a function call, but “outside” the process
 - Does not have the address of the system function to call
 - Like a Remote Procedure Call (RPC) – for later
 - Marshall the syscall id and args in registers and exec syscall
- Interrupt
 - External asynchronous event triggers context switch
 - eg. Timer, I/O device
 - Independent of user process
- Trap or Exception
 - Internal synchronous event in process triggers context switch
 - e.g., Protection violation (segmentation fault), Divide by zero, ...

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.30

So how do we safely let user hand over a task to kernel ?

Implementing Safe Kernel Mode Transfers

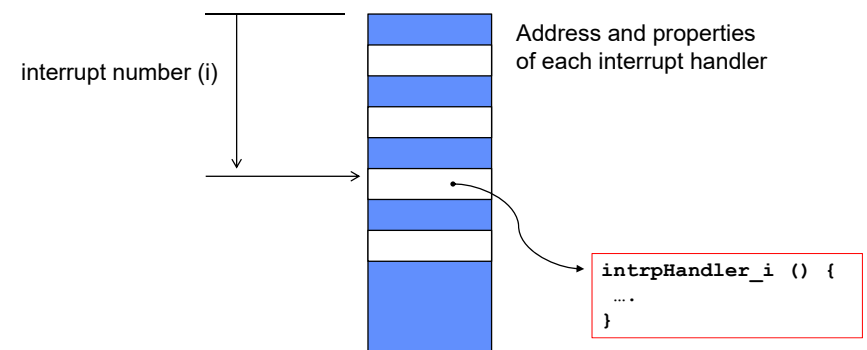
- Important aspects:
 - Controlled transfer into kernel (e.g., syscall table)
 - Separate kernel stack
 - Carefully constructed kernel code packs up the user process state and sets it aside
 - Details depend on the machine architecture
- Trust NOTHING !**
- Should be impossible for buggy or malicious user program to cause the kernel to corrupt itself

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.31

Interrupt Vector



- Where else do you see this dispatch pattern?

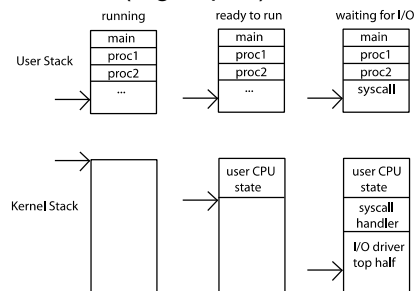
1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.32

Need for Separate Kernel Stacks

- Kernel needs space to work
- Cannot put anything on the user stack (Why?) 2 threads, 1 entering kernel mode
- Two-stack model
 - OS thread has interrupt stack (located in kernel memory) plus User stack (located in user memory)
 - Syscall handler copies user args to kernel space before invoking specific function (e.g., open)
 - Interrupts (???)

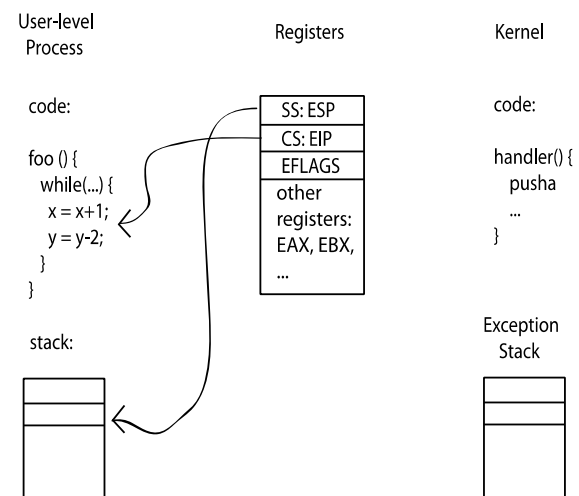


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.33

Before

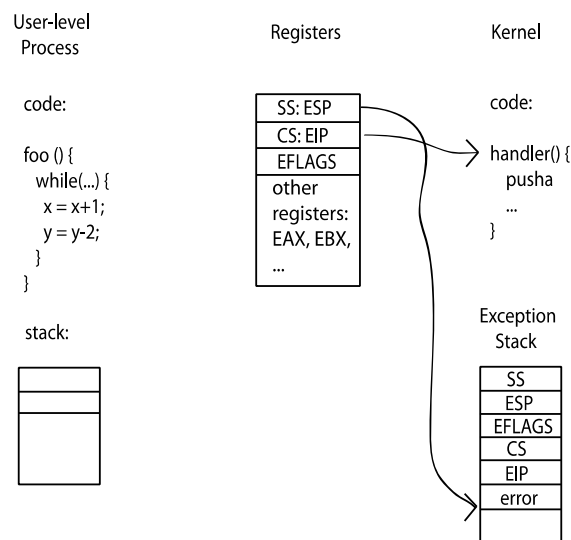


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.34

During During The Interrupt



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.35

Kernel System Call Handler

- **Vector through well-defined syscall entry points!**
 - Table mapping system call number to handler
- **Locate arguments**
 - In registers or on user (!) stack
- **Copy arguments**
 - From user memory into kernel memory
 - Protect kernel from malicious code evading checks
- **Validate arguments**
 - Protect kernel from errors in user code
- **Copy results back**
 - Into user memory

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.36

Hardware support: Interrupt Control

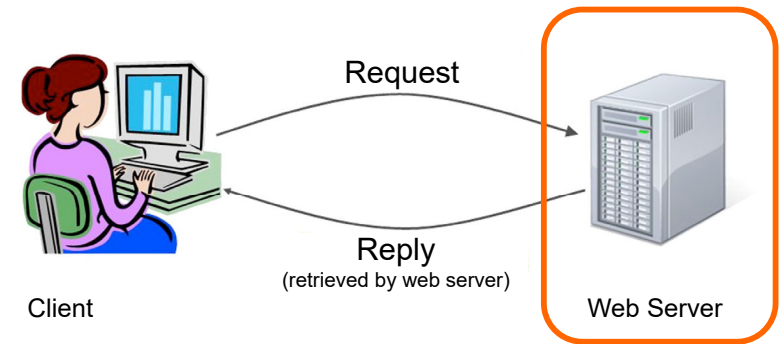
- Interrupt processing not visible to the user process:
 - Occurs between instructions, restarted transparently
 - No change to process state
 - What can be observed even with perfect interrupt processing?
- Interrupt Handler invoked with interrupts 'disabled'
 - Re-enabled upon completion
 - Non-blocking (run to completion, no waits)
 - Pack up in a queue and pass off to an OS thread for hard work
 - » wake up an existing OS thread

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.37

Putting it together: web server

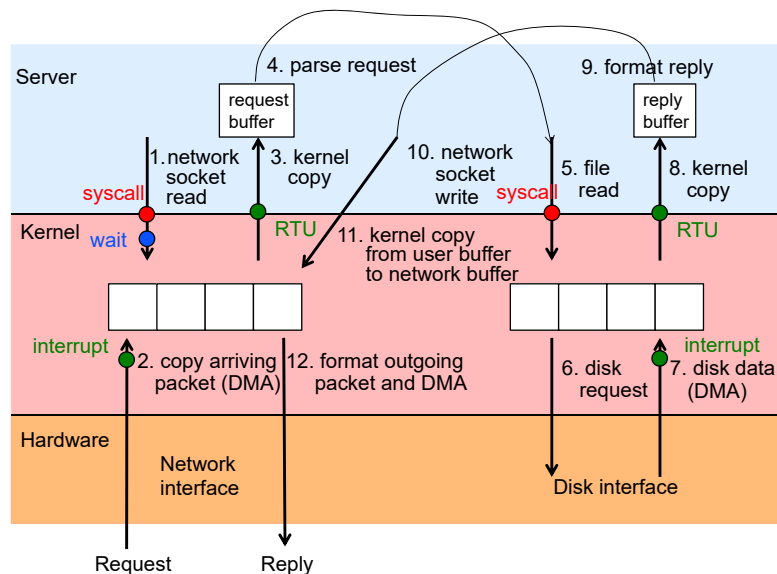


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.38

Putting it together: web server



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.39

A DEEP Dive into Pintos
and
Where we are Going!

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.40

Project 1: Processes

- Allocate and initialize Process object
- Allocate and initialize kernel thread mini-stack and associated Thread object
- Allocate and initialize page table for process
- Load code and static data into user pages
- Build initial User Stack
 - Initial register contents
- Schedule (post) process thread for execution
- ...
- Eventually switch to user thread ...
- Several lists of various types

Project 1

Pintos: `process.c`, `thread.c`

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.41

Understanding "Address Space"

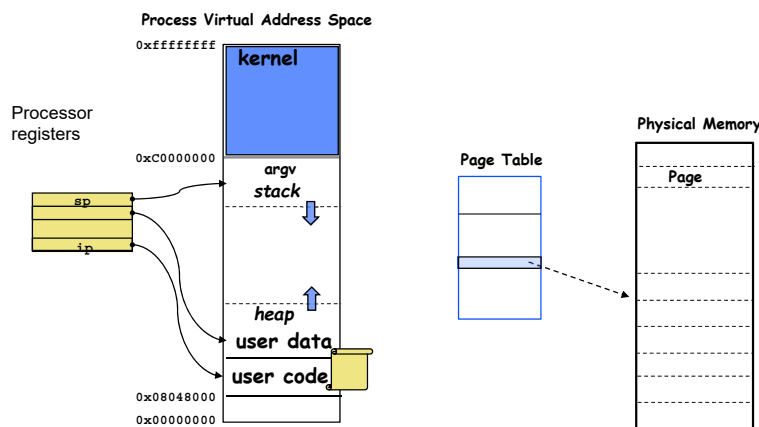
- Page table is the primary mechanism
- Privilege Level determines which regions can be accessed
 - Which entries can be used
- System (PL=0) can access all, User (PL=3) only part
- Each process has its own address space
- The "System" part of all of them is the same
 - ⇒ All system threads share the same system address space and same memory
- This address pattern less (not?) common now after the Meltdown attack was discovered in 2017
 - More Later in Term!!

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.42

User Process View

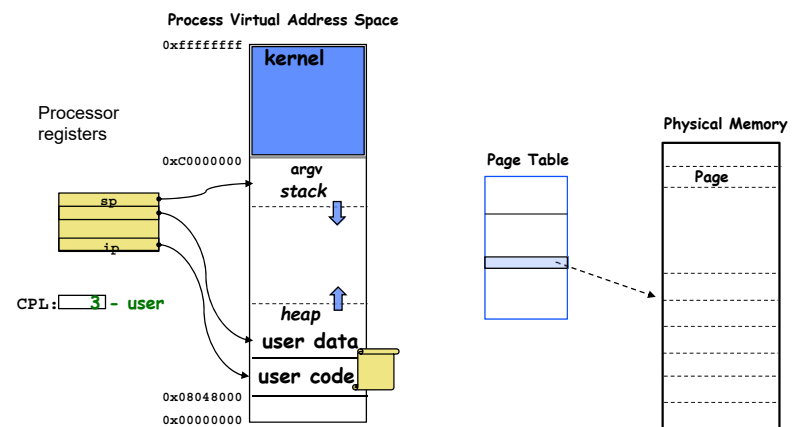


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.43

Processor Mode (Privilege Level)

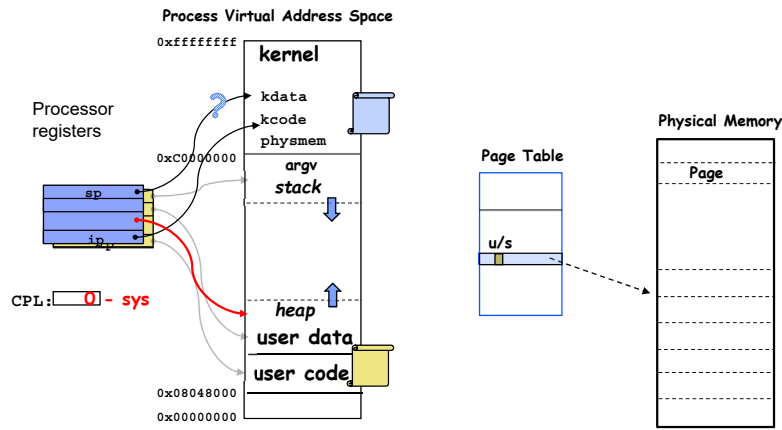


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.44

User → Kernel: PL = 0

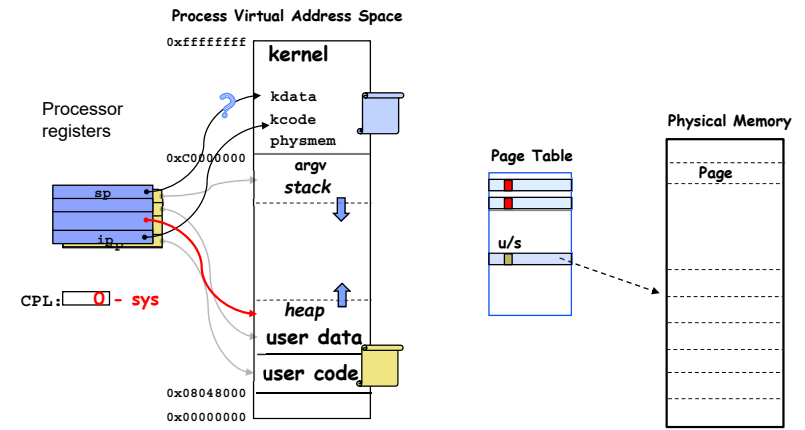


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.45

Page Table enforces PL

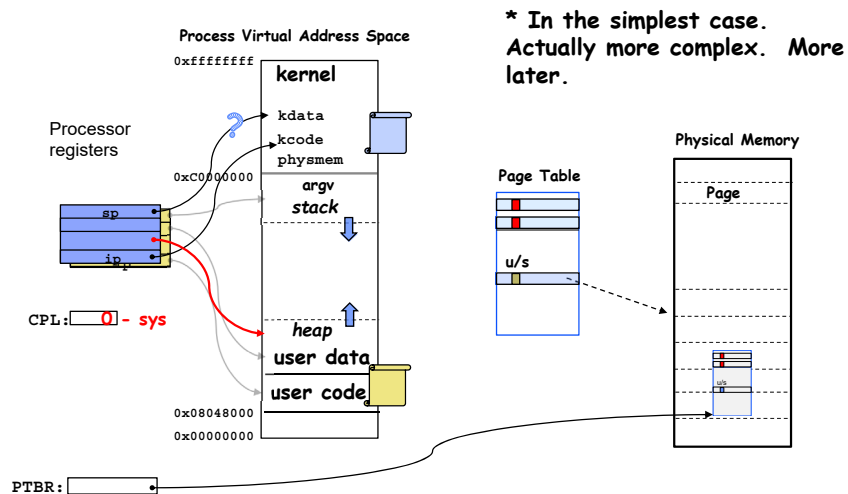


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.46

Page Table resides in memory*

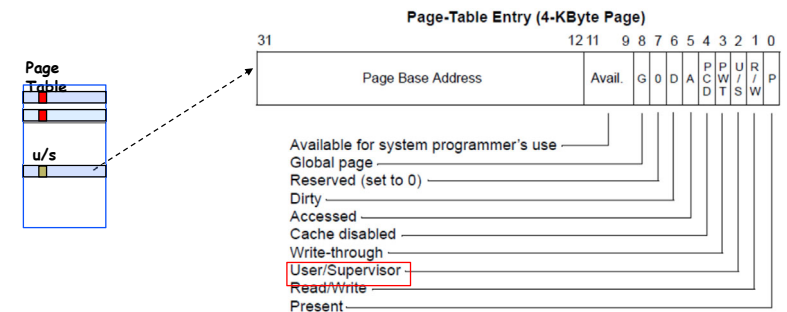


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.47

x86 (32-bit) Page Table Entry



- Controls many aspects of access
- Later – discuss page table organization
 - For 32 (64?) bit VAS, how large? vs size of memory?
 - Use sparsely. Very very fast HW access

Pintos: [page_dir.c](#)

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.48

Kernel Portion of Address Space

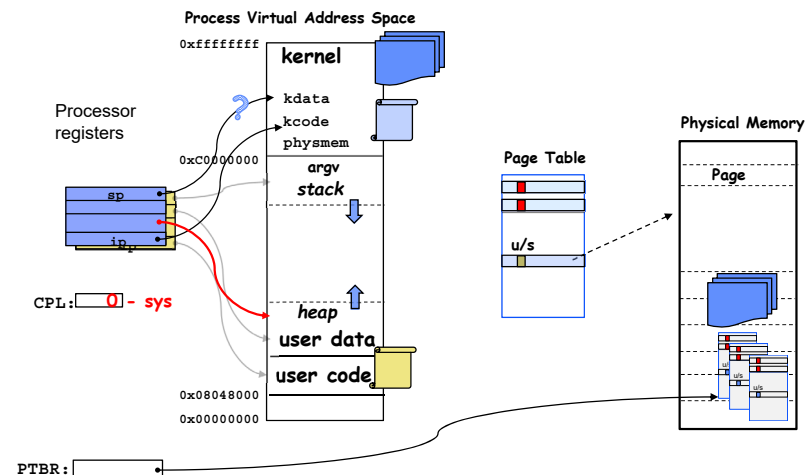
- Contains the kernel code
 - Loaded when the machine booted
- Explicitly mapped to physical memory
 - OS creates the page table
- Used to contain all kernel data structures
 - List of all the processes and threads
 - The page tables for those processes
 - Other system resources (files, sockets, ttys, ...)
- Also contains (little) stacks for “kernel threads”
 - Early OS design serviced all processes on a single execution thread
 - » Event driven programming
 - Today: Each Process Thread supported by (little) Kernel Thread

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.49

1 Kernel Code, many Kernel “stacks”



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.50

From Machine Structure to OS Data Structures

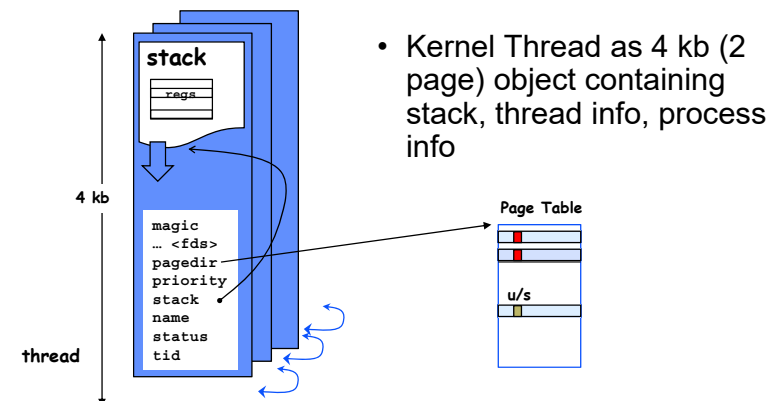
- Traditional Unix, etc. design maintains a Process Control Block (PCB) per process
- Each with a Thread Control Block (TCB) per thread of that process
- Today, assume **single** thread per process
 - PINTOS model
- Linux organized around threads with “groups of threads” associated with a process

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.51

PINTOS Thread



- Kernel Thread as 4 kb (2 page) object containing stack, thread info, process info

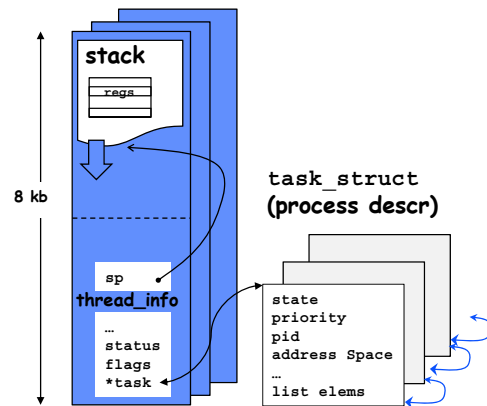
Pintos: thread.c

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.52

Linux “Task”



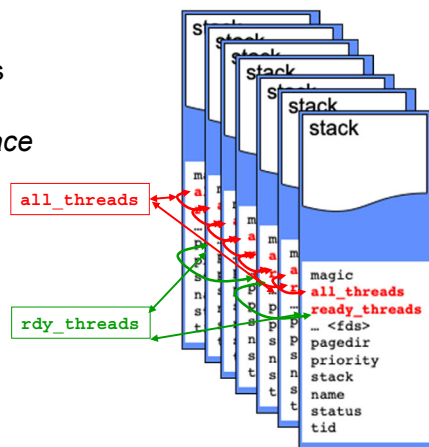
- Kernel Thread as 8 kb (2 page) object containing stack and thread information + process descriptor

Process Creation

- Allocate and initialize Process object
- Allocate and initialize kernel thread mini-stack and associated Thread object
- Allocate and initialize page table for process
 - Referenced by process object
- Load code and static data into user pages
- Build initial User Stack
 - Initial register contents, argv, ...
- Schedule (post) process thread for execution
- ...
- Eventually *switch* to user thread ...
- Several lists of various types

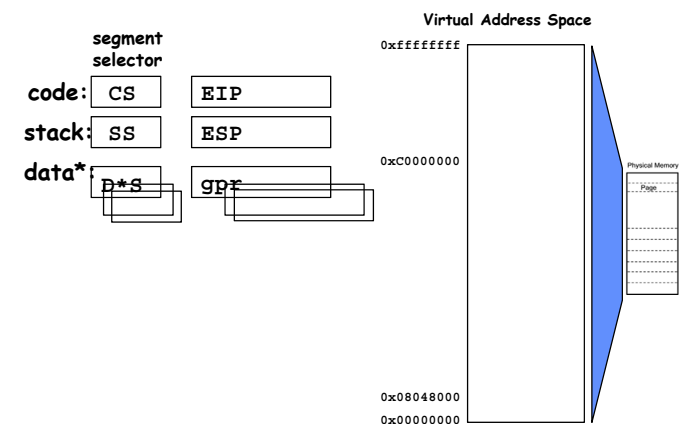
Aside: Polymorphic lists in C

- Many places in the kernel need to maintain a “list of X”
 - This is tricky in C, which has no polymorphism
 - Essentially adding an *interface* to a package (ala Go)
- In Linux and Pintos this is done by embedding a `list_elem` in the struct
 - Macros allow shift of view between object and list
 - You’ll practice in HW1 – before getting into PINTOS

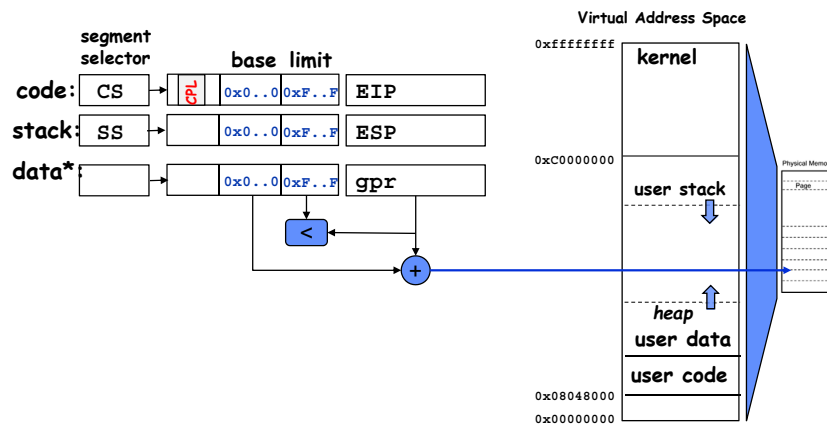


Pintos: list.c

Bit of x86 thread/process/VAS management



Bit of x86 thread/process/VAS management

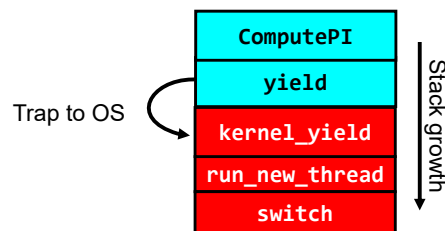


Pintos: loader.h

Recall: 3 types of U→K Mode Transfer

- Syscall
 - Process requests a system service, e.g., exit
 - Like a function call, but “outside” the process
 - Does not have the address of the system function to call
 - Like a Remote Procedure Call (RPC) – for later
 - Marshall the syscall id and args in registers and exec syscall
- Interrupt
 - External asynchronous event triggers context switch
 - eg. Timer, I/O device
 - Independent of user process
- Trap
 - Internal synchronous event in process triggers context switch
 - e.g., Protection violation (segmentation fault), Divide by zero, ...
- All 3 *exceptions* are an UNPROGRAMMED CONTROL TRANSFER
 - Where does it go? (To handler specified in interrupt vector)
 - Are interrupts enabled or disabled when get there?

Stack for Thread Transition

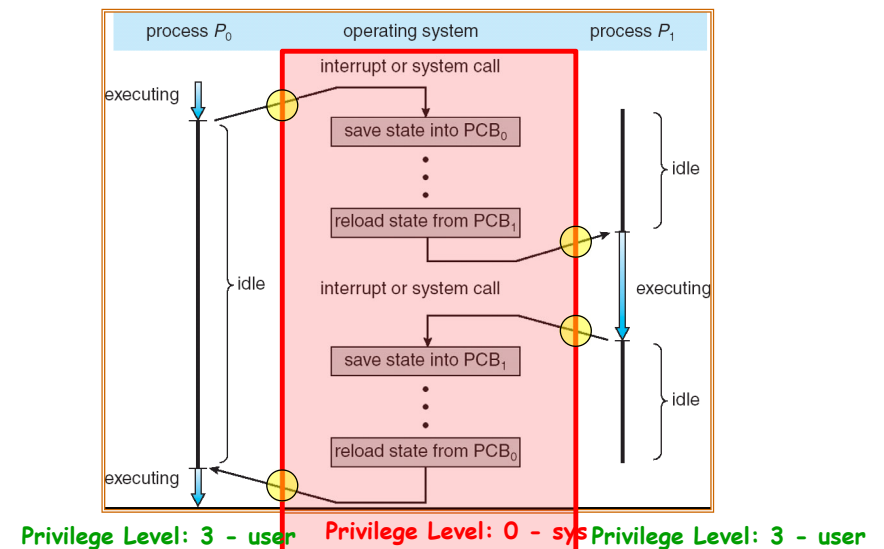


Cyan = User Stack; Red = Kernel Stack

```
run_new_thread() {
    newThread = PickNewThread();
    switch(curThread, newThread)
    ThreadHouseKeeping(); /* Do any cleanup */
}
```

Scheduling: Policy Decision

A Privileged View of the Context Switch



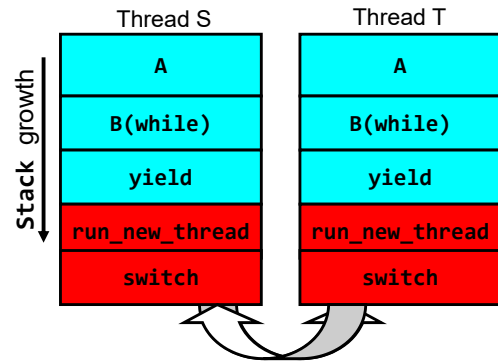
Stacks During Context Switch

- Consider the following code blocks:

```

proc A() {
    B();
}
proc B() {
    while(TRUE) {
        yield();
    }
}
    
```

- Suppose we have 2 threads:
 - Threads S and T



Saving/Restoring state (often called "Context Switch")

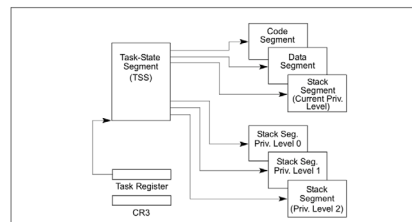
```

Switch(tCur, tNew) {
    /* Unload old thread */
    TCB[tCur].regs.r7 = CPU.r7;
    ...
    TCB[tCur].regs.r0 = CPU.r0;
    TCB[tCur].regs.sp = CPU.sp;
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/

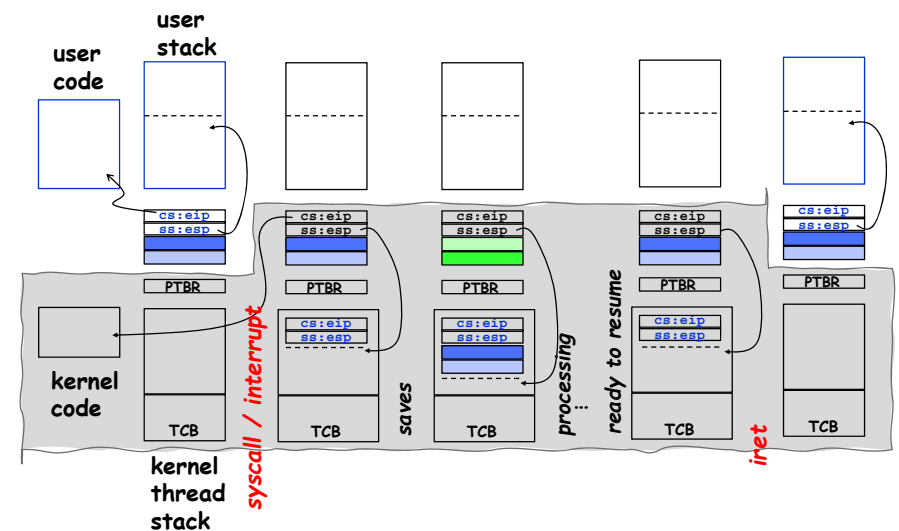
    /* Load and execute new thread */
    CPU.r7 = TCB[tNew].regs.r7;
    ...
    CPU.r0 = TCB[tNew].regs.r0;
    CPU.sp = TCB[tNew].regs.sp;
    CPU.retpc = TCB[tNew].regs.retpc;
    return; /* Return to CPU.retpc */
}
    
```

Hardware context switch support

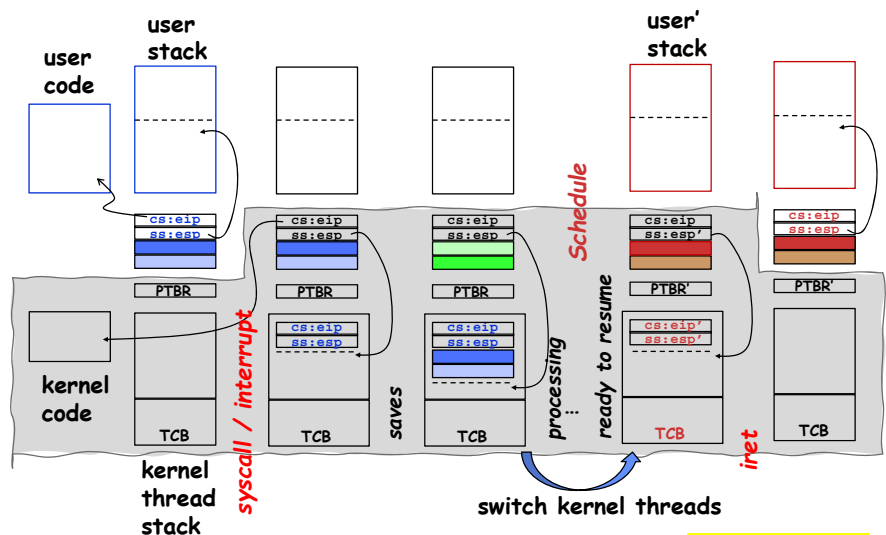
- Syscall/Intr (U → K)
 - PL 3 → 0;
 - TSS ← EFLAGS, CS:EIP;
 - SS:SP ← k-thread stack (TSS PL 0);
 - push (old) SS:ESP onto (new) k-stack
 - push (old) eflags, cs:eip, <err>
 - CS:EIP ← <k target handler>
- Then
 - Handler then saves other regs, etc
 - Does all its work, possibly choosing other threads, changing PTBR (CR3)
- iret (K → U)
 - PL 0 → 3;
 - Eflags, CS:EIP ← popped off k-stack
 - SS:SP ← user thread stack (TSS PL 3);



Context Switch – in pictures



Context Switch – Scheduling



Pintos: switch.S

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.65

Concurrency

- But, ... ???
- With all these threads in the kernel, won't they step on each other?
 - For example, while one is loading a program, other threads should run ...
 - Processes are isolated from each other, but all the threads in the kernel share the kernel address space, memory, data structures
- We will study synchronization soon
- The kernel controls whether hardware interrupts are enabled or not
 - Disabled on entry, selectively enable
 - Atomic operations, ...

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.66

Dispatch Loop

```

Loop {
    RunThread();
    ChooseNextThread();
    SaveStateOfCPU(curTCB);
    LoadStateOfCPU(newTCB);
}
    
```

- Conceptually all the OS executes
- Infinite Loop
 - When would we ever "exit?"
 - Can we assume some thread is always ready?

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.67

Dispatch Loop

```

Loop {
    RunThread();
    ChooseNextThread();
    SaveStateOfCPU(curTCB);
    LoadStateOfCPU(newTCB);
}
    
```

How to run a new thread?

- Load thread's registers into CPU
- Load its environment (address space, if in different process)
- Jump to thread's PC

How does dispatch loop get control again?

- Thread returns control voluntarily – yield, I/O
- External events: thread is preempted

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.68

Thread Operations in Pintos

- `thread_create(name, priority, func, args)`
 - Create a new thread to run `func(args)`
- `thread_yield()`
 - Relinquish processor voluntarily
- `thread_join(thread)`
 - Wait (put in queue) until thread exits, then return
- `thread_exit`
 - Quit thread and clean up, wake up joiner if any

*More later,
incl. synch
ops*

Meta-Question

- Process is an instance of a program executing.
 - The fundamental OS responsibility
- Processes do their work by processing and calling file system operations
- Are there any operations on processes themselves?
- `exit` ?

pid.c

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>
int main(int argc, char *argv[])
{
    pid_t pid = getpid();    /* get current processes PID */

    printf("My pid: %d\n", pid);

    exit(0);
}
```

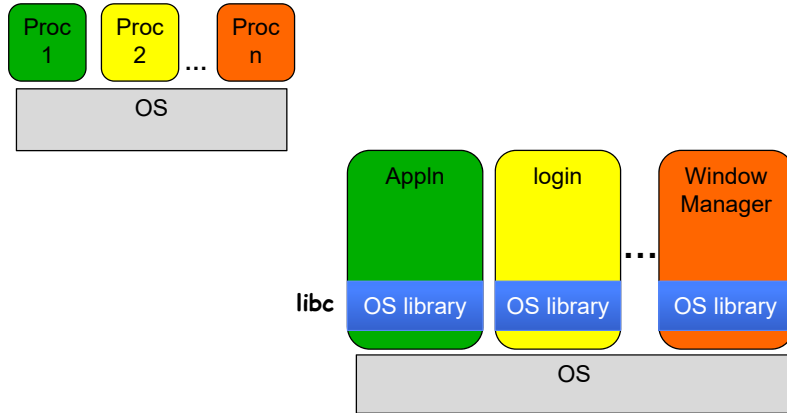
ps anyone?

Can a process create a process ?

- Yes
- Fork creates a copy of process
- What about the program you want to run?

see Lec 3.80

OS Run-Time Library

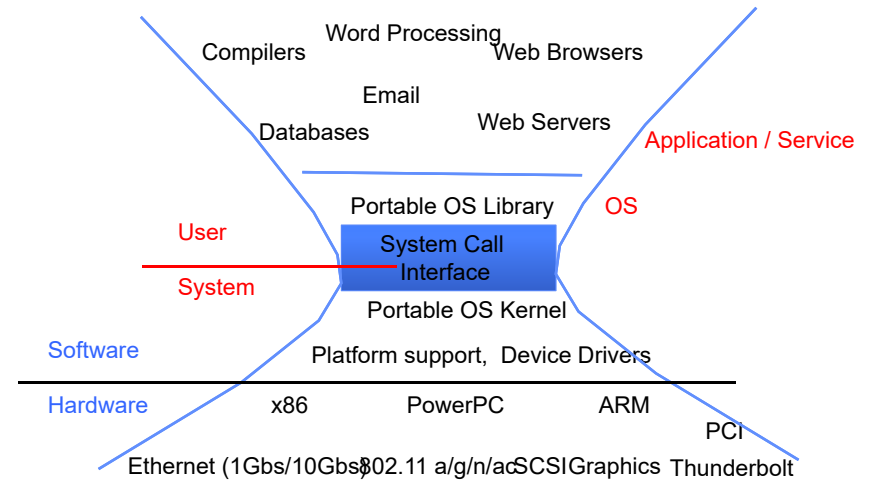


1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.73

A Narrow Waist



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.74

POSIX/Unix

- **Portable Operating System Interface [X?]**
- Defines “Unix”, derived from AT&T Unix
 - Created to bring order to many Unix-derived OSs
- Interface for **application programmers** (mostly)

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.75

System Calls

Application:

```
fd = open(pathname);
```

Library:

```
File *open(pathname) {
    asm code ... syscall # into ax
    put args into registers bx, ...
    special trap instruction
```

Operating System:

```
get args from regs
dispatch to system func
process, schedule, ...
complete, resume process
```

```
get results from regs
```

```
};
```

Continue with results

Pintos: userprog/syscall.c, lib/user/syscall.c

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.76

SYSCALLs (of over 300)

%eax	Name	Source	%ebx	%ecx	%edx	%esi	%edi
1	sys_exit	kernel/exit.c	int	-	-	-	-
2	sys_fork	arch/i386/kernel/process.c	struct pt_regs	-	-	-	-
3	sys_read	fs/read_write.c	unsigned int	char *	size_t	-	-
4	sys_write	fs/read_write.c	unsigned int	const char *	size_t	-	-
5	sys_open	fs/open.c	const char *	int	int	-	-
6	sys_close	fs/open.c	unsigned int	-	-	-	-
7	sys_waitpid	kernel/exit.c	pid_t	unsigned int *	int	-	-
8	sys_creat	fs/open.c	const char *	int	-	-	-
9	sys_link	fs/namei.c	const char *	const char *	-	-	-
10	sys_unlink	fs/namei.c	const char *	-	-	-	-
11	sys_execve	arch/i386/kernel/process.c	struct pt_regs	-	-	-	-
12	sys_chdir	fs/open.c	const char *	-	-	-	-
13	sys_time	kernel/time.c	int *	-	-	-	-
14	sys_mknod	fs/namei.c	const char *	int	dev_t	-	-
15	sys_chmod	fs/open.c	const char *	mode_t	-	-	-
16	sys_lchown	fs/open.c	const char *	uid_t	gid_t	-	-
18	sys_stat	fs/stat.c	char *	struct _old_kernel_stat *	-	-	-
19	sys_lseek	fs/read_write.c	unsigned int	off_t	unsigned int	-	-
20	sys_getpid	kernel/sched.c	-	-	-	-	-
21	sys_mount	fs/super.c	char *	char *	char *	-	-
22	sys_oldumount	fs/super.c	char *	-	-	-	-
23	sys_setuid	kernel/sys.c	uid_t	-	-	-	-
24	sys_getuid	kernel/sched.c	-	-	-	-	-
25	sys_stime	kernel/time.c	int *	-	-	-	-
26	sys_ptrace	arch/i386/kernel/ptrace.c	long	long	long	long	-
27	sys_alarm	kernel/sched.c	unsigned int	-	-	-	-
28	sys_fstat	fs/stat.c	unsigned int	struct _old_kernel_stat *	-	-	-
29	sys_pause	arch/i386/kernel/sys_i386.c	-	-	-	-	-
30	sys_ftime	fs/open.c	char *	struct utimbuf *	-	-	-

Pintos: syscall-nr.h

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.77

Recall: Kernel System Call Handler

- Locate arguments
 - In registers or on user(!) stack
- Copy arguments
 - From user memory into kernel memory
 - Protect kernel from malicious code evading checks
- Validate arguments
 - Protect kernel from errors in user code
- Copy results back
 - into user memory

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.78

Process Management

- exit – terminate a process
- **fork – copy the current process**
- exec – change the *program* being run by the current process
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.79

Creating Processes

- pid_t fork(); -- copy the current process
 - New process has different pid
- Return value from **fork()**: pid (like an integer)
 - When > 0:
 - » Running in (original) **Parent** process
 - » return value is **pid** of new child
 - When = 0:
 - » Running in new **Child** process
 - When < 0:
 - » Error! Must handle somehow
 - » Running in original process
- **State of original process duplicated in *both* Parent and Child!**
 - Address Space (Memory), File Descriptors (covered later), etc...

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.80

fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.81

fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.82

fork1.c

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.83

fork_race.c

```
int i;
cpid = fork();
if (cpid > 0) {
    for (i = 0; i < 10; i++) {
        printf("Parent: %d\n", i);
        // sleep(1);
    }
} else if (cpid == 0) {
    for (i = 0; i > -10; i--) {
        printf("Child: %d\n", i);
        // sleep(1);
    }
}
```

- What does this print?
- Would adding the calls to sleep matter?

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.84

Fork “race”

```
int i;
cpid = fork();
if (cpid > 0) {
    for (i = 0; i < 10; i++) {
        printf("Parent: %d\n", i);
        // sleep(1);
    }
} else if (cpid == 0) {
    for (i = 0; i > -10; i--) {
        printf("Child: %d\n", i);
        // sleep(1);
    }
}
```



Process Management

- fork – copy the current process
- exec – change the *program* being run by the current process
- **wait – wait for a process to finish**
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

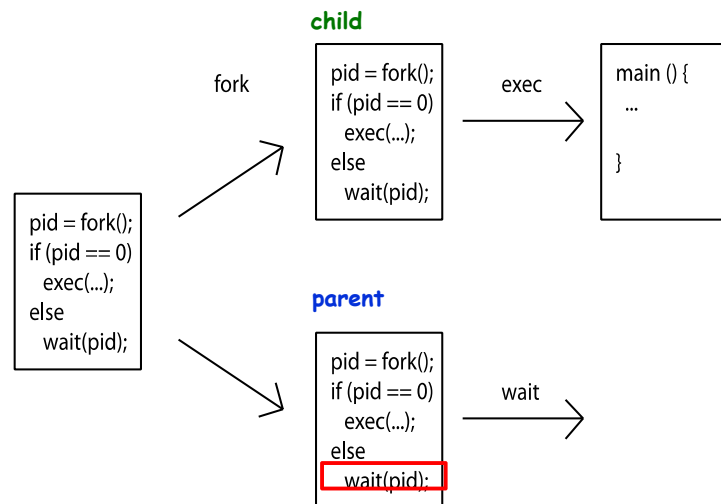
fork2.c – parent waits for child to finish

```
int status;
pid_t tcpid;
...
cpid = fork();
if (cpid > 0) {                /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
    tcpid = wait(&status);
    printf("[%d] bye %d(%d)\n", mypid, tcpid, status);
} else if (cpid == 0) {        /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
}
...
```

Process Management

- fork – copy the current process
- **exec – change the *program* being run by the current process**
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- sigaction – set handlers for signals

Process Management



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.89

fork3.c

```
...
cpid = fork();
if (cpid > 0) { /* Parent Process */
    tcpid = wait(&status);
} else if (cpid == 0) { /* Child Process */
    char *args[] = {"ls", "-l", NULL};
    execv("/bin/ls", args);
    /* execv doesn't return when it works.
       So, if we got here, it failed! */
    perror("execv");
    exit(1);
}
...
```

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.90

Shell

- A shell is a job control system
 - Allows programmer to create and manage a set of programs to do some task
 - Windows, MacOS, Linux all have shells
- Example: to compile a C program

```
cc -c sourcefile1.c
cc -c sourcefile2.c
ln -o program sourcefile1.o sourcefile2.o
./program
```



1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.91

Process Management

- fork – copy the current process
- exec – change the *program* being run by the current process
- wait – wait for a process to finish
- kill – send a *signal* (interrupt-like notification) to another process
- **sigaction – set handlers for signals**

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.92

inf_loop.c

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <signal.h>

void signal_callback_handler(int signum) {
    printf("Caught signal!\n");
    exit(1);
}

int main() {
    struct sigaction sa;
    sa.sa_flags = 0;
    sigemptyset(&sa.sa_mask);
    sa.sa_handler = signal_callback_handler;

    sigaction(SIGINT, &sa, NULL);
    while (1) {}
}
```

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.93

Common POSIX Signals

- SIGINT – control-C
- SIGTERM – default for kill shell command
- SIGSTP – control-Z (default action: stop process)

- SIGKILL, SIGSTOP – terminate/stop process
 - Can't be changed or disabled with sigaction
 - Why?

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.94

Summary

- Process consists of two pieces
 1. Address Space (Memory & Protection)
 2. One or more threads (Concurrency)
- Represented in kernel as
 - Process object (resources associated with process)
 - Thread object + (mini) stack
 - Hardware support critical in U → K → U context switch
 - Different privileges in different modes (CPL, Page Table)
- Variety of process management syscalls
 - fork, exec, wait, kill, sigaction
- Scheduling: Threads move between queues
- Threads: multiple stacks per address space
 - Context switch: Save/Restore registers, "return" from new thread's switch routine
 - So far, we've only seen kernel threads

1/28/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 3.95