

CS162

Operating Systems and Systems Programming

Lecture 5

Concurrency and Mutual Exclusion

February 4th, 2020
Prof. John Kubiatowicz
<http://cs162.eecs.Berkeley.edu>

Recall: Fork, Wait, and (optional) Exec

```
cpid = fork();
if (cpid > 0) { // Parent Process
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
    tcpid = wait(&status);
    printf("[%d] bye %d\n", mypid, tcpid);
} else if (cpid == 0) { // Child Process
    mypid = getpid();
    printf("[%d] child\n", mypid);
    execl(filename, (char *)0); // Opt: start new program
} else { // Error! }
```

- Return value from Fork: integer
 - When > 0: return value is pid of new child (Running in **Parent**)
 - When = 0: Running in new **Child** process
 - When < 0: Error! Must handle somehow
- Wait() system call: wait for next child to exit
 - Return value is PID of terminating child
 - Argument is pointer to integer variable to hold exit status
- Exec() family of calls: replace process with new executable

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.2

The way how shell works ?

Recall: Internal Events

- Blocking on I/O
 - The act of requesting I/O implicitly yields the CPU
- Waiting on a “signal” from other thread
 - Thread asks to wait and thus yields the CPU
- Thread executes a yield()
 - Thread volunteers to give up CPU


```
computePI() {
    while(TRUE) {
        ComputeNextDigit();
        yield();
    }
}
```

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.3

Recall: Stack for Yielding Thread



- How do we run a new thread?

```
run_new_thread() {
    newThread = PickNewThread();
    switch(curThread, newThread);
    ThreadHouseKeeping(); /* Do any cleanup */
}
```

- How does dispatcher switch to a new thread?
 - Save anything next thread may trash: PC, regs, stack pointer
 - Maintain isolation for each thread

2/4/20


Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.4

Recall: Multithreaded Stack Switching

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```




- Suppose we have 2 threads:
 - Threads S and T

Goals for Today

- Finish discussion of Threads
- Concurrency and need for Synchronization Operations
- Basic Synchronization through Locks
- Initial Lock Implementations



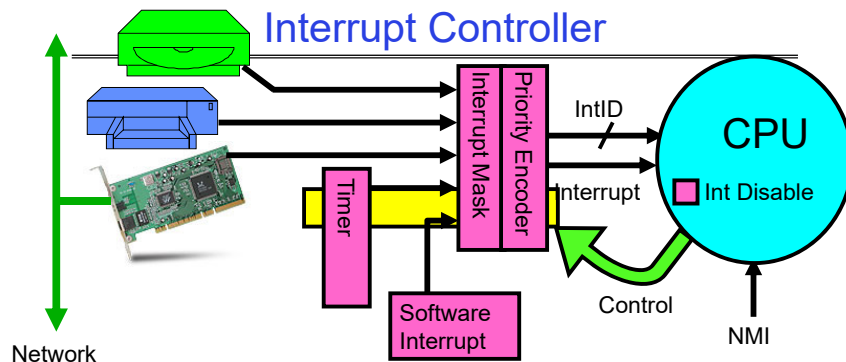
What happens when thread blocks on I/O?



- What happens when a thread requests a block of data from the file system?
 - User code invokes a system call
 - Read operation is initiated
 - Run new thread/switch
- Thread communication similar
 - Wait for Signal/Join
 - Networking

External Events

- What happens if thread never does any I/O, never waits, and never yields control?
 - Could the ComputePI program grab all resources and never release the processor?
 - » What if it didn't print to console?
 - Must find way that dispatcher can regain control!
- Answer: utilize external events
 - Interrupts: signals from hardware or software that stop the running code and jump to kernel
 - Timer: like an alarm clock that goes off every some milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs



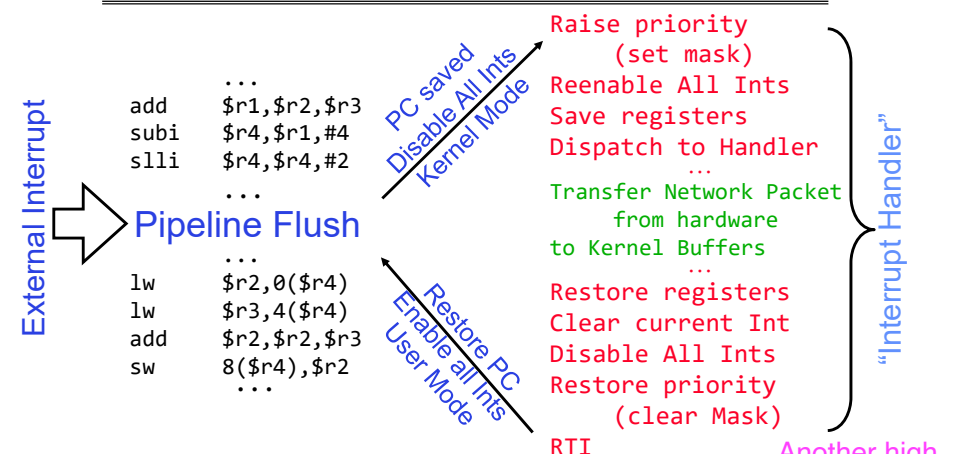
- Interrupts invoked with interrupt lines from devices
- Interrupt controller chooses interrupt request to honor
 - Interrupt identity specified with ID line
 - Mask enables/disables interrupts
 - Priority encoder picks highest enabled interrupt
 - Software Interrupt Set/Cleared by Software
- CPU can disable all interrupts with internal flag
- Non-Maskable Interrupt line (NMI) can't be disabled

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.9

Example: Network Interrupt



- An interrupt is a hardware-invoked context switch
 - No separate step to choose what to run next
 - Always run the interrupt handler immediately

Another high priority interrupt may come in when an interrupt is handling.


2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.10

Use of Timer Interrupt to Return Control

- Solution to our dispatcher problem
 - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {
    DoPeriodicHouseKeeping();
    run_new_thread();
}
```

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.11

Hardware context switch support in x86

- Syscall/Intr (U → K)
 - PL 3 → 0;
 - TSS ← EFLAGS, CS:EIP;
 - SS:SP ← k-thread stack (TSS PL 0);
 - push (old) SS:ESP onto (new) k-stack
 - push (old) eflags, cs:eip, <err>
 - CS:EIP ← <k target handler>

task segment

level goes 3 to 0




Figure 7-1. Structure of a Task

- Then
 - Handler then saves other regs, etc
 - Does all its work, possibly choosing other threads, changing PTBR (CR3)
- kernel thread has set up user GPRs
- iret (K → U)
 - PL 0 → 3;
 - Eflags, CS:EIP ← popped off k-stack
 - SS:SP ← user thread stack (TSS PL 3);

pg 2,942 of 4,922 of x86 reference manual

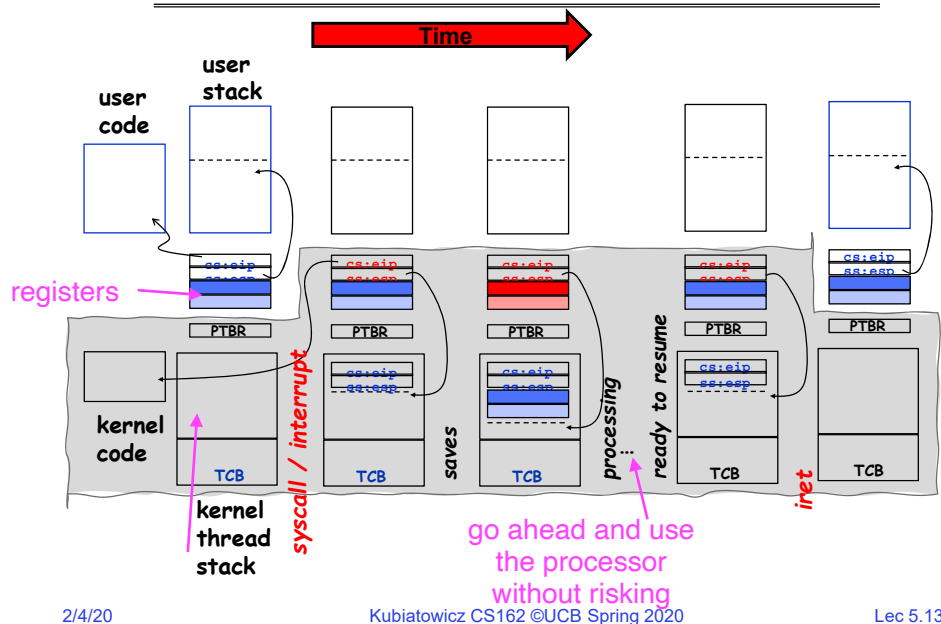
Pintos: tss.c, intr-stubs.S

2/4/20

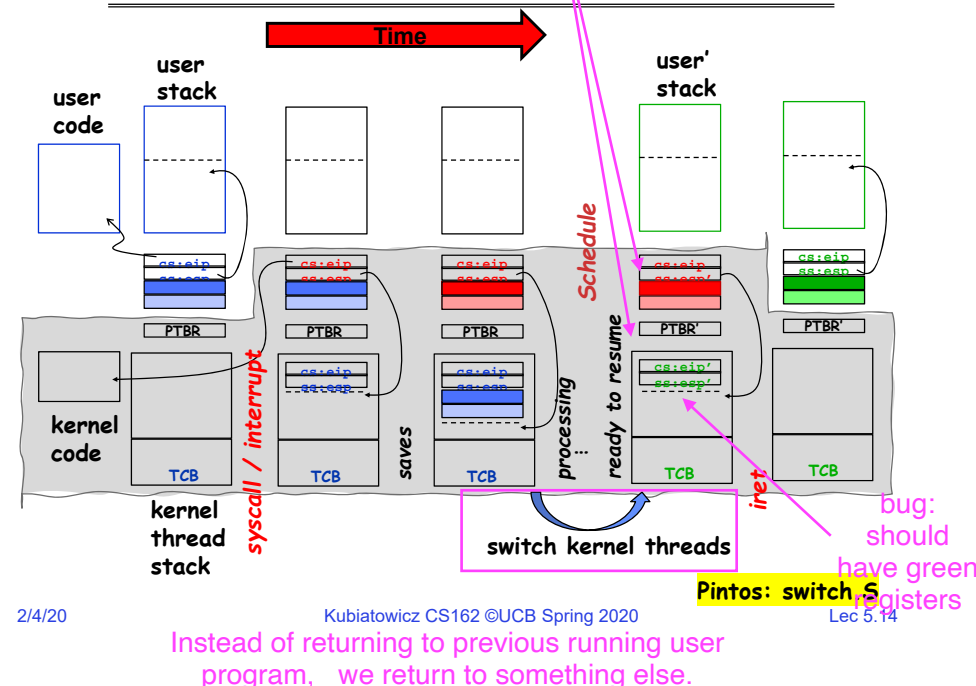
Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.12

Pintos: Kernel Crossing on Syscall or Interrupt



Pintos: Context Switch – Scheduling



ThreadFork(): Create a New Thread

- ThreadFork() is a user-level procedure that creates a new thread and places it on ready queue
- Arguments to ThreadFork()
 - Pointer to application routine (fcnPtr)
 - Pointer to array of arguments (fcnArgPtr)
 - Size of stack to allocate
- Implementation
 - Sanity check arguments
 - Enter Kernel-mode and Sanity Check arguments again
 - Allocate new Stack and TCB
 - Initialize TCB and place on ready list (Runnable)

How do we initialize TCB and Stack?


- Initialize Register fields of TCB
 - Stack pointer made to point at stack
 - PC return address \Rightarrow OS (asm) routine ThreadRoot()
 - Two arg registers (say rdi and rsi for x86) initialized to fcnPtr and fcnArgPtr, respectively
- Initialize stack data?
 - No. Important part of stack frame is in registers (ra)
 - Think of stack frame as just before body of ThreadRoot() really gets started

ThreadRoot stub

Stack growth

Initial Stack

How does Thread get started?



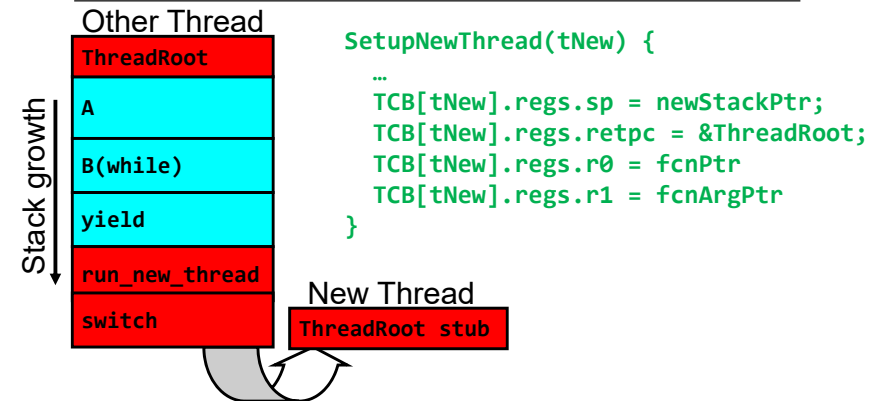
- Need to construct a new kernel thread that is ready to run when switch goes to it
- Note that switch doesn't know any difference between new or preexisting thread!

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.17

How does a thread get started?



- How do we make a **new** thread?
 - Setup TCB/kernel thread to point at new user stack and ThreadRoot code
 - Put pointers to start function and args in registers
 - This depends heavily on the calling convention (i.e. RISC-V vs x86)
- Eventually, run_new_thread() will select this TCB and return into beginning of ThreadRoot()
- This really starts the new thread

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.18


What does ThreadRoot() look like?

- ThreadRoot() is the root for the thread routine:

```

ThreadRoot(fcnPTR, fcnArgPtr) {
    DoStartupHousekeeping();
    UserModeSwitch(); /* enter user mode */
    Call fcnPtr(fcnArgPtr);
    ThreadFinish();
}
    
```

- Startup Housekeeping
 - Includes things like recording start time of thread
 - Other statistics
- Stack will grow and shrink with execution of thread
- Final return from thread returns into ThreadRoot() which calls ThreadFinish()
 - ThreadFinish() wake up sleeping threads



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.19

Administrivia

- Group Creation Deadline is TONIGHT!
 - Need 4 people in a group
 - If you signup with less, we may end up adding another person to your group!
- All members of a group need to have the same TA
 - Priority for same section; if cannot make this work, keep same TA
 - Remember: Your TA needs to see you in section!

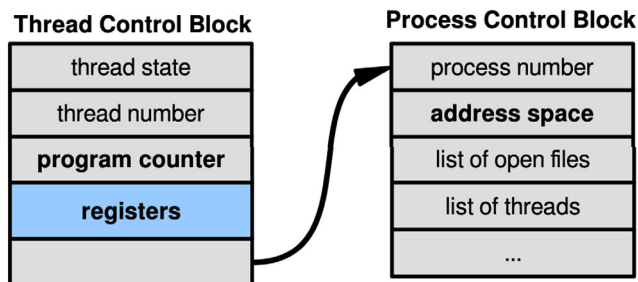
2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.20

Kernel-Supported Threads

- Each thread has a **thread control block**
 - CPU registers, including PC, pointer to stack
 - Scheduling info: priority, etc.
 - Pointer to **Process control block**
- OS scheduler uses TCBs, not PCBs

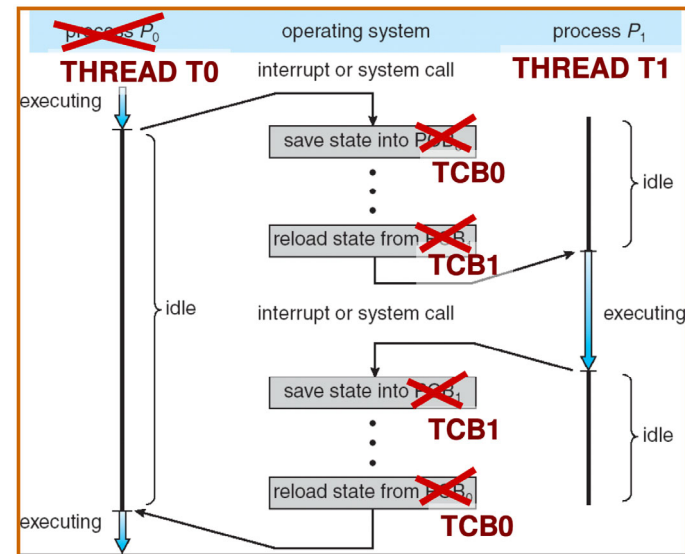


2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.21

Kernel-Supported User Threads



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.22

User-level Multithreading: *pthread*s

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine)(void*), void *arg);`
 - thread is created executing *start_routine* with *arg* as its sole argument. (return is implicit call to `pthread_exit`)
- `void pthread_exit(void *value_ptr);`
 - terminates and makes *value_ptr* available to any successful join
- `int pthread_join(pthread_t thread, void **value_ptr);`
 - suspends execution of the calling thread until the target *thread* terminates.
 - On return with a non-NULL *value_ptr* the value passed to `pthread_exit()` by the terminating thread is made available in the location referenced by *value_ptr*.

man pthread
<https://pubs.opengroup.org/onlinepubs/7908799/xsh/pthread.h.html>

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.23

Little Example

How to tell if something is done?
 Really done?
 OK to reclaim its resources?

```

(base) CullerMac19:code04 culler$ ./pthread 4
Main stack: 7ffee2c6b6b8, common: 10cf95048 (162)
Thread #1 stack: 70000d83bef8 common: 10cf95048 (162)
Thread #3 stack: 70000d941ef8 common: 10cf95048 (164)
Thread #2 stack: 70000d8beef8 common: 10cf95048 (165)
Thread #0 stack: 70000d7b8ef8 common: 10cf95048 (163)

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <string.h>

int common = 162;

void *threadfun(void *threadid)
{
    long tid = (long)threadid;
    printf("Thread #%lx stack: %lx common: %lx (%d)\n", tid,
        (unsigned long) &tid, (unsigned long) &common, common++);
    pthread_exit(NULL);
}


int main (int argc, char *argv[])
{
    long t;
    int nthreads = 2;
    if (argc > 1) {
        nthreads = atoi(argv[1]);
    }
    pthread_t *threads = malloc(nthreads*sizeof(pthread_t));
    printf("Main stack: %lx, common: %lx (%d)\n",
        (unsigned long) &t, (unsigned long) &common, common);
    for(t=0; t<nthreads; t++){
        pthread_create(&threads[t], NULL, threadfun, (void *)t);
        if (rc) {
            printf("ERROR: return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    for(t=0; t<nthreads; t++){
        pthread_join(threads[t], NULL);
    }
    pthread_exit(NULL);
}
    
```

2/4/20

Lec 5.24

Fork-Join Pattern




- Main thread *creates* (forks) collection of sub-threads passing them args to work on, *joins* with them, collecting results.

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.25

Thread Abstraction




- Illusion: Infinite number of processors

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.26

Thread Abstraction



- Illusion: Infinite number of processors
- Reality: Threads execute with variable speed
 - Programs must be designed to work with any schedule

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.27

Programmer vs. Processor View

Programmer's View	Possible Execution #1
.	.
.	.
.	.
x = x + 1;	x = x + 1;
y = y + x;	y = y + x;
z = x + 5y;	z = x + 5y;
.	.
.	.
.	.

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.28

Programmer vs. Processor View

Programmer's View	Possible Execution #1	Possible Execution #2
.	.	.
.	.	.
.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1$
$y = y + x;$	$y = y + x;$
$z = x + 5y;$	$z = x + 5y;$	thread is suspended
.	.	other thread(s) run
.	.	thread is resumed
.
		$y = y + x$
		$z = x + 5y$

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.29

Programmer vs. Processor View



Programmer's View	Possible Execution #1	Possible Execution #2	Possible Execution #3
.	.	.	.
.	.	.	.
.	.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1$	$x = x + 1$
$y = y + x;$	$y = y + x;$	$y = y + x$
$z = x + 5y;$	$z = x + 5y;$	thread is suspended
.	.	other thread(s) run	thread is suspended
.	.	thread is resumed	other thread(s) run
.	thread is resumed
		$y = y + x$
		$z = x + 5y$	$z = x + 5y$

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.30

Possible Executions



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.31

Per Thread Descriptor (Kernel Supported Threads)

- Each Thread has a **Thread Control Block (TCB)**
 - Execution State: CPU registers, program counter (PC), pointer to stack (SP)
 - Scheduling info: state, priority, CPU time
 - Various Pointers (for implementing scheduling queues)
 - Pointer to enclosing process (PCB) – user threads
 - ... (add stuff as you find a need)
- OS Keeps track of TCBs in “kernel memory”
 - In Array, or Linked List, or ...
 - I/O state (file descriptors, network connections, etc)


2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.32

Multithreaded Processes

- Process Control Block (PCBs) points to multiple Thread Control Blocks (TCBs):



- Switching threads within a block is a simple thread switch
- Switching threads across blocks requires changes to memory and I/O address tables


2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.33

Multiprocessing vs Multiprogramming

- Remember Definitions:
 - Multiprocessing \equiv Multiple CPUs
 - Multiprogramming \equiv Multiple Jobs or Processes
 - Multithreading \equiv Multiple threads per Process
- What does it mean to run two threads “concurrently”?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.34

Correctness for systems with concurrent threads

- If dispatcher can schedule threads in any way, programs must work under all circumstances
 - Can you test for this?
 - How can you know if your program works?
- Independent Threads:**
 - No state shared with other threads
 - Deterministic \Rightarrow Input state determines results
 - Reproducible \Rightarrow Can recreate Starting Conditions, I/O
 - Scheduling order doesn't matter (if `switch()` works!!!)
- Cooperating Threads:**
 - Shared State between multiple threads
 - Non-deterministic
 - Non-reproducible
- Non-deterministic and Non-reproducible means that bugs can be intermittent
 - Sometimes called “Heisenbugs”

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.35

Interactions Complicate Debugging

- Is any program truly independent?
 - Every process shares the file system, OS resources, network, etc
 - Extreme example: buggy device driver causes thread A to crash “independent thread” B
- You probably don't realize how much you depend on reproducibility:
 - Example: Evil C compiler
 - Modifies files behind your back by inserting errors into C program unless you insert debugging code
 - Example: Debugging statements can overrun stack
- Non-deterministic errors are really difficult to find
 - Example: Memory layout of kernel+user programs
 - depends on scheduling, which depends on timer/other things
 - Original UNIX had a bunch of non-deterministic errors
 - Example: Something which does interesting I/O
 - User typing of letters used to help generate secure keys

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.36

Why allow cooperating threads?


- People cooperate; computers help/enhance people's lives, so computers must cooperate
 - By analogy, the non-reproducibility/non-determinism of people is a notable problem for “carefully laid plans”
- Advantage 1: Share resources
 - One computer, many users
 - One bank balance, many ATMs
 - » What if ATMs were only updated at night?
 - Embedded systems (robot control: coordinate arm & hand)
- Advantage 2: Speedup
 - Overlap I/O and computation
 - » Many different file systems do read-ahead
 - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
 - More important than you might think
 - Chop large problem up into simpler pieces
 - » To compile, for instance, gcc calls `cpp | cc1 | cc2 | as | ld`
 - » Makes system easier to extend

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.37

High-level Example: Web Server



- Server must handle many requests
- Non-cooperating version:

```
serverLoop() {
    con = AcceptCon();
    ProcessFork(ServiceWebPage(), con);
}
```
- What are some disadvantages of this technique?

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.38

Threaded Web Server

- Now, use a single process
- Multithreaded (cooperating) version:

```
serverLoop() {
    connection = AcceptCon();
    ThreadFork(ServiceWebPage(), connection);
}
```
- Looks almost the same, but has many advantages:
 - Can share file caches kept in memory, results of CGI scripts, other things
 - Threads are *much* cheaper to create than processes, so this has a lower per-request overhead
- Question: would a user-level (say one-to-many) thread package make sense here?
 - When one request blocks on disk, all block...
- What about Denial of Service attacks or digg / Slash-dot effects?

2/4/20


Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.39



Thread Pools

- Problem with previous version: Unbounded Threads
 - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```
master() {
    allocThreads(worker, queue);
    while(TRUE) {
        con=AcceptCon();
        Enqueue(queue, con);
        wakeUp(queue);
    }
}
```


```
worker(queue) {
    while(TRUE) {
        con=Dequeue(queue);
        if (con==null)
            sleepOn(queue);
        else
            ServiceWebPage(con);
    }
}
```

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.40

ATM Bank Server



- ATM server problem:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.41

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```
BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- How could we speed this up?
 - More than one request being processed at once
 - Event driven (overlap computation and I/O)
 - Multiple threads (multi-proc, or overlap comp and I/O)

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.42

Event Driven Version of ATM server

- Suppose we only had one CPU
 - Still like to overlap I/O with computation
 - Without threads, we would have to rewrite in event-driven style
- Example

```
BankServer() {
    while(TRUE) {
        event = WaitForNextEvent();
        if (event == ATMRequest)
            StartOnRequest();
        else if (event == AcctAvail)
            ContinueRequest();
        else if (event == AcctStored)
            FinishRequest();
    }
}
```

- What if we missed a blocking I/O step?
- What if we have to split code into hundreds of pieces which could be blocking?
- This technique is used for graphical programming

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.43

Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without “deconstructing” code into non-blocking fragments
 - One thread per request
- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* May use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- Unfortunately, shared state can get corrupted:

<u>Thread 1</u>	<u>Thread 2</u>
load r1, acct->balance	
	load r1, acct->balance
	add r1, amount2
	store r1, acct->balance
add r1, amount1	
store r1, acct->balance	

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.44

Problem is at the Lowest Level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

x = 1;

Thread B

y = 2;

- However, what about (Initially, y = 12):

Thread A

x = 1;

x = y+1;

Thread B

y = 2;

y = y*2;

- What are the possible values of x?

- Or, what are the possible values of x below?

Thread A

x = 1;

Thread B

x = 2;

- X could be 1 or 2 (non-deterministic!)

- Could even be 3 for serial processors:

» Thread A writes 0001, B writes 0010 → scheduling order
ABABABBA yields 3!

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.45

Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- Atomic Operation:** an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
 - Consequently – weird example that produces “3” on previous slide can't happen
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to copy a whole array

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.46

Another Concurrent Program Example

- Two threads, A and B, compete with each other
 - One tries to increment a shared counter
 - The other tries to decrement the counter

Thread A

i = 0;

while (i < 10)

i = i + 1;

printf("A wins!");

Thread B

i = 0;

while (i > -10)

i = i - 1;

printf("B wins!");

- Assume that memory loads and stores are atomic, but incrementing and decrementing are *not* atomic
- Who wins? Could be either
- Is it guaranteed that someone wins? Why or why not?
- What if both threads have their own CPU running at same speed? Is it guaranteed that it goes on forever?

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.47

Hand Simulation Multiprocessor Example

- Inner loop looks like this:

Thread A

r1=0 load r1, M[i]

r1=1 add r1, r1, 1

M[i]=1 store r1, M[i]

Thread B

r1=0 load r1, M[i]

r1=-1 sub r1, r1, 1

M[i]=-1 store r1, M[i]

- Hand Simulation:**
 - And we're off. A gets off to an early start
 - B says “hmp, better go fast” and tries really hard
 - A goes ahead and writes “1”
 - B goes and writes “-1”
 - A says “HUH??? I could have sworn I put a 1 there”
- Could this happen on a uniprocessor? With Hyperthreads?
 - Yes! Unlikely, but if you are depending on it not happening, it will and your system will break...

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.48

Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
 - Cooperating threads inherently non-deterministic and non-reproducible
 - Really hard to debug unless carefully designed!
- Example: Therac-25
 - Machine for radiation therapy
 - Software control of electron accelerator and electron beam/Xray production
 - Software control of dosage
 - Software errors caused the death of several patients
 - A series of race conditions on shared variables and poor software design
 - "They determined that data entry speed during editing was the key factor in producing the error condition: If the prescription data was edited at a fast pace, the overdose occurred."




Figure 1. Typical Therac-25 facility

Motivating Example: "Too Much Milk"

- Great thing about OS's – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

Definitions

- Synchronization**: using atomic operations to ensure cooperation between threads
 - For now, only loads and stores are atomic
 - We are going to show that its hard to build anything useful with only reads and writes
- Mutual Exclusion**: ensuring that only one thread does a particular thing at a time
 - One thread *excludes* the other while doing its task
- Critical Section**: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
 - Critical section is the result of mutual exclusion
 - Critical section and mutual exclusion are two ways of describing the same thing

More Definitions

- Lock**: prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much: roommate angry if only wants OJ



- Of Course – We don't know how to make a lock yet

Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Impulse is to start coding first, then when it doesn't work, pull hair out
 - Instead, think first, then code
 - Always write down behavior first
- What are the correctness properties for the "Too much milk" problem???
- Never more than one person buys
- Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

2/4/20

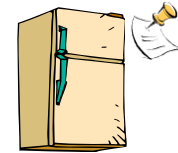
Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.53

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of "lock")
 - Remove note after buying (kind of "unlock")
 - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.54

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of "lock")
 - Remove note after buying (kind of "unlock")
 - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

Thread A

```
if (noMilk) {  
  
    if (noNote) {  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```

Thread B

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```

2/4/20


Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.55

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of "lock")
 - Remove note after buying (kind of "unlock")
 - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



- Result?
 - Still too much milk **but only occasionally!**
 - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the dispatcher does!

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.56

Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;
if (noMilk) {
    if (noNote) {
        buy milk;
    }
}
remove Note;
```

- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk



Too Much Milk Solution #2


- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

```
Thread A
leave note A;
if (noNote B) {
    if (noMilk) {
        buy Milk;
    }
}
remove note A;
```

```
Thread B
leave note B;
if (noNoteA) {
    if (noMilk) {
        buy Milk;
    }
}
remove note B;
```

- Does this work?
- Possible for neither thread to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
 - Extremely unlikely this would happen, but will at worse possible time
 - Probably something like this in UNIX

Too Much Milk Solution #2: problem!



- I'm not getting milk, You're getting milk
- This kind of lockup is called "starvation!"

Too Much Milk Solution #3

- Here is a possible two-note solution:

```
Thread A
leave note A;
while (note B) {\\X
    do nothing;
}
if (noMilk) {
    buy milk;
}
remove note A;
```

```
Thread B
leave note B;
if (noNote A) {\\Y
    if (noMilk) {
        buy milk;
    }
}
remove note B;
```

- Does this work? **Yes**. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - If no note B, safe for A to buy,
 - Otherwise wait to find out what will happen
- At Y:
 - If no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

Case 1

- “leave note A” happens before “if (noNote A)”

```

leave note A;
while (note B) {\X
  do nothing;
};

if (noMilk) {
  buy milk;}
remove note A;

leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
  
```

Diagram: A blue arrow labeled "happened before" points from the "leave note A;" line to the "if (noNote A) {\Y" line.

Case 1

- “leave note A” happens before “if (noNote A)”

```

leave note A;
while (note B) {\X
  do nothing;
};

if (noMilk) {
  buy milk;}
remove note A;

leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
  
```

Diagram: A blue arrow labeled "happened before" points from the "leave note A;" line to the "if (noNote A) {\Y" line.

Case 1

- “leave note A” happens before “if (noNote A)”

```

leave note A;
while (note B) {\X
  do nothing;
};

if (noMilk) {
  buy milk;}
remove note A;

leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
  
```

Diagram: A blue arrow labeled "happened before" points from the "leave note A;" line to the "if (noNote A) {\Y" line. A dashed arrow points from the "remove note B;" line to the "if (noMilk) {" line. A vertical dashed line with the text "Wait for note B to be removed" is positioned between the "leave note A;" block and the "if (noMilk) {" block.

Case 2

- “if (noNote A)” happens before “leave note A”

```

leave note A;
while (note B) {\X
  do nothing;
};


if (noMilk) {
  buy milk;}
remove note A;

leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
  
```

Diagram: A blue arrow labeled "happened before" points from the "if (noNote A) {\Y" line to the "leave note A;" line.

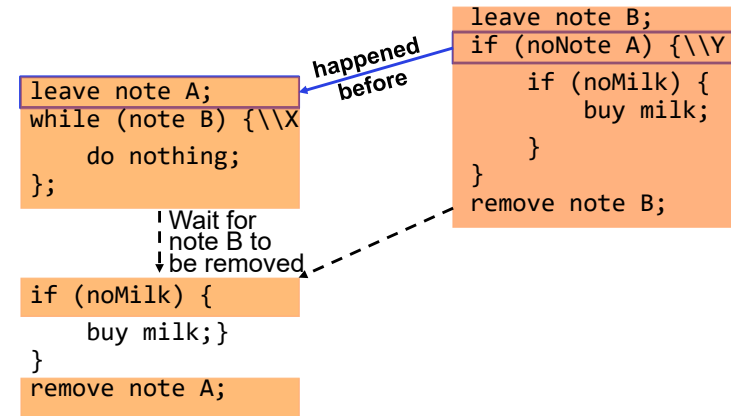
Case 2

- “if (noNote A)” happens before “leave note A”



Case 2

- “if (noNote A)” happens before “leave note A”



Solution #3 discussion

- Our solution protects a single “Critical-Section” piece of code for each thread:

```

if (noMilk) {
    buy milk;
}
    
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - » Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - » Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - » This is called “busy-waiting”
- There's a better way
 - Have hardware provide higher-level primitives than atomic load & store
 - Build even higher-level programming abstractions on this hardware support

Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock
 - **lock.Acquire()** – wait until lock is free, then grab
 - **lock.Release()** – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:


```

milklock.Acquire();
if (nomilk)
    buy milk;
milklock.Release();
            
```
- Once again, section of code between Acquire() and Release() called a “**Critical Section**”
- Of course, you can make this even simpler: suppose you are out of ice cream instead of milk
 - Skip the test since you always need more ice cream ;-)

How to Implement Locks?

- **Lock:** prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting
 - » Should *sleep* if waiting for a long time
- Atomic Load/Store: get solution like Milk #3
 - Pretty complex and error prone
- Hardware Lock instruction
 - Is this a good idea?
 - What about putting a task to sleep?
 - » What is the interface between the hardware and scheduler?
 - Complexity?
 - » Done in the Intel 432
 - » Each feature makes HW more complex and slow



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020


Lec 5.69

Naïve use of Interrupt Enable/Disable

- How can we build multi-instruction atomic operations?
 - Recall: dispatcher gets control in two ways.
 - » Internal: Thread does something to relinquish the CPU
 - » External: Interrupts cause dispatcher to take CPU
 - On a uniprocessor, can avoid context-switching by:
 - » Avoiding internal events (although virtual memory tricky)
 - » Preventing external events by disabling interrupts
- Consequently, naïve Implementation of locks:


```
LockAcquire { disable Ints; }
LockRelease { enable Ints; }
```
- Problems with this approach:
 - **Can't let user do this!** Consider following:


```
LockAcquire();
While(TRUE) {;
```
 - Real-Time system—no guarantees on timing!
 - » Critical Sections might be arbitrarily long
 - What happens with I/O or other important events?
 - » “Reactor about to meltdown. Help?”



2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.70

Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;
```



```
Acquire() {
  disable interrupts;
  if (value == BUSY) {
    put thread on wait queue;
    Go to sleep();
    // Enable interrupts?
  } else {
    value = BUSY;
  }
  enable interrupts;
}

Release() {
  disable interrupts;
  if (anyone on wait queue) {
    take thread off wait queue;
    Place on ready queue;
  } else {
    value = FREE;
  }
  enable interrupts;
}
```

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.71

Where are we going with synchronization?

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Send/Receive
Hardware	Load/Store Disable Ints Test&Set Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
 - Everything is pretty painful if only atomic primitives are load and store
 - Need to provide primitives useful at user-level

2/4/20

Kubiatowicz CS162 ©UCB Spring 2020

Lec 5.72

Summary

- Concurrent threads are a very useful abstraction
 - Allow transparent overlapping of computation and I/O
 - Allow use of parallel processing when available
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleavings
 - Without careful design, shared variables can become completely inconsistent
- Important concept: Atomic Operations
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives