

Section 1: OS Concepts, Processes, Threads

CS 162

January 31, 2020

Contents

1	Vocabulary	2
2	Fundamental Operating System Concepts	4
3	Processes	6
3.1	Forks	6
3.2	Process Stack Allocation	6
3.3	Process Heap Allocation	7
3.4	Simple Wait	7
3.5	Exec	8
3.6	Exec + Fork	8
4	Threads	9
4.1	Join	9
4.2	Thread Stack Allocation	10
4.3	Thread Heap Allocation	10
4.4	The Central Galactic Floopy Corporation	11
5	Interrupt Handlers	13
5.1	Pintos Interrupt Handler	14

1 Vocabulary

With credit to the Anderson & Dahlin textbook (A&D):

- **process** - A process is an instance of a computer program that is being executed, typically with restricted rights. It consists of an address space and one or more threads of control. It is the main abstraction for protection provided by the operating system kernel.
- **thread** - A thread is a single execution sequence that can be managed independently by the operating system. (See A&D, 4.2)
- **isolation** - Isolating (separating) applications from one another so that a potentially misbehaving application cannot corrupt other applications or the operating system.
- **dual-mode operation** - Dual-mode operation refers to hardware support for multiple privilege levels: a privileged level (called *supervisor-mode* or *kernel-mode*) that provides unrestricted access to the hardware, and a restricted level (called *user-mode*) that executes code with restricted rights.
- **privileged instruction** - Instruction available in kernel mode but not in user mode. Two examples of privileged instructions are the instructions to enable and disable interrupts on the processor. If user-level code could disable interrupts, it would guarantee that the user-level process could run on a hardware thread for as long as it wanted.
- **unprivileged instruction** - Instruction available in both user mode and kernel mode. An example of an unprivileged instruction is the **add** instruction or the instructions that read or write to memory. User-level processes are allowed to perform these standard operations that all computer programs need in order to run.
- **fork** - A C function that calls the fork syscall that creates a new process by duplicating the calling process. The new process, referred to as the child, is an exact duplicate of the calling process (except for a few details, read more in the man page). Both the newly created process and the parent process return from the call to fork. On success, the PID of the child process is returned in the parent, and 0 is returned in the child. On failure, -1 is returned in the parent, no child process is created.
- **wait** - A class of C functions that call syscalls that are used to wait for state changes in a child of the calling process, and obtain information about the child whose state has changed. A state change is considered to be: the child terminated; the child was stopped by a signal; or the child was resumed by a signal.
- **exit code** - The exit status or return code of a process is a 1 byte number passed from a child process (or callee) to a parent process (or caller) when it has finished executing a specific procedure or delegated task
- **exec** - The `exec()` family of functions replaces the current process image with a new process image. The initial argument for these functions is the name of a file that is to be executed.
- **pthread** - A POSIX-compliant (standard specified by IEEE) implementation of threads. A `pthread_t` is usually just an alias for “unsigned long int”.
- **pthread_create** - Creates and immediately starts a child thread running in the same address space of the thread that spawned it. The child executes starting from the function specified. Internally,

this is implemented by calling the clone syscall.

```
/* On success, pthread_create() returns 0; on error, it returns an error
 * number, and the contents of *thread are undefined. */
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine) (void *), void *arg);
```

- **pthread_join** - Waits for a specific thread to terminate, similar to waitpid(3).

```
/* On success, pthread_join() returns 0; on error, it returns an error number. */
int pthread_join(pthread_t thread, void **retval);
```

- **critical section** - A section of code that accesses a shared resource and must not be concurrently run by more than a single thread.
- **race condition** - A situation whose outcome is dependent on the sequence of execution of multiple threads running simultaneously.
- **lock** - Synchronization primitives that provide mutual exclusion. Threads may acquire or release a lock. Only one thread may hold a lock at a time. If a thread attempts to acquire a lock that is held by some other thread, it will block at that line of code until the lock is released and it successfully acquires it. Implementations can vary.

2 Fundamental Operating System Concepts

1. What are the 3 roles the OS plays?

Referee: The OS manages the sharing of resources, isolation and protect from other processes.
Illusionist: The OS provides a clean and easy to use abstraction of the underlying hardware resources.
Glue: The OS provides common services between processes to facilitate sharing, community, and user level uniformity.

2. How is a process different from a thread?

A thread is an independent execution context, with its own registers, program counter, and stack. A thread is defined by an OS object called the Thread Control Block (TCB) which stores all of this information.

A process is comprised of one or more threads and enjoys OS level isolation from other operation systems, this includes having its own address space. The process is defined by an OS object called the Process Control Block (PCB) which stores process level information like a pointer to a page table, a list of open files, and process metadata.

Different threads in the same process may each have their own stacks, but all share the process' address space and so can access each other's memory. **Threads encapsulate concurrency (modern OS schedule by threads not processes) while processes encapsulate isolation.**

3. What is the process address space and address translation? Why are they important?

A process works with virtual memory and address translation is used to map the process virtual memory to the machine's physical memory. This is important because

- (a) Virtual memory provides an isolation abstraction that gives the illusion of the process being the sole user of the address space.
- (b) Gives the illusion of a large address space without having to actually allocate that much physical memory.
- (c) Provides isolation between processes because virtual memory between processes do not (usually) translate to the same physical memory, so different processes will not be able to access each other's memory.

4. What is dual mode operation and what are the three forms of control transfer from user to kernel mode?

Dual mode helps provide isolation between processes by restricting the use of certain resources and actions to just the system in "kernel mode."

The three forms of transfer are:

- (1) If a process wants to use those resources, like the filesystem, it can perform a control transfer called **a system call or syscall** from user space to kernel space. The kernel then validates the arguments of the syscall and accesses the resource on the behalf of the process.
- (2) **An interrupt** is an external asynchronous event independent of the user process that requires the attention of the operating system. The kernel switches from the process to

the kernel in order to respond to this interrupt, which can be something like an OS timer going off or incoming bits on an I/O device.

- (3) **A trap** is an internal synchronous event such as a fault in the execution of the user process that forces the process to transfer control to the kernel. It is caused by instructions in the user process such as dividing by zero or accessing invalid memory.

5. Why does a thread in kernel mode have a separate kernel stack? What can happen if the kernel stack was in the user address space?

Each OS thread has a kernel stack (located in kernel memory) plus an user stack (located in user memory). This kernel stack is memory that is only accessible by the thread in privileged kernel mode and thus provides protection for the kernel from the user.

If the user process had write permission to the kernel stack, it could corrupt the kernel in the event that the user purposely or accidentally overwrite the kernel stack. Some OSes keep the PCB in the kernel stack, which if the user had access to can allow for permissions changes, information leakage, and breaking of process isolation.

6. How does the syscall handler protect the kernel from corrupt or malicious user code?

- The syscall number is used to index into a vector mapping number of fixed syscall handlers — this prevents the user from getting the kernel to run user code in kernel mode.
- The handler copies the user arguments from the user registers or stack onto the kernel stack — this protects the kernel from malicious code on the user stack from evading checks.
- All of the arguments are validated before the syscall is executed — this protects the kernel from errors in the user arguments like invalid or null pointers.
- Results from the syscall are copied back into user memory so the caller has access to them.

3 Processes

3.1 Forks

How many new processes are created in the below program assuming calls to fork succeeds?

```
int main(void)
{
    for (int i = 0; i < 3; i++)
        pid_t pid = fork();
}
```

7 (8 including the original process). Newly forked processes will continue to execute the loop from wherever their parent process left off.

3.2 Process Stack Allocation

What can C print?

```
int main(void)
{
    int stuff = 5;
    pid_t pid = fork();
    printf("The last digit of pi is %d\n", stuff);
    if (pid == 0)
        stuff = 6;
}
```

The last digit of pi is 5
The last digit of pi is 5
(Since the entire address space is copied, stuff will still remain the same. There is also the fork() failure case where only one line is printed)

3.3 Process Heap Allocation

What can C print?

```
int main(void)
{
    int* stuff = malloc(sizeof(int)*1);
    *stuff = 5;
    pid_t pid = fork();
    printf("The last digit of pi is %d\n", *stuff);
    if (pid == 0)
        *stuff = 6
}
```

The last digit of pi is 5
The last digit of pi is 5
(Since the entire address space is copied, stuff will still remain the same. There is also the fork() failure case where only one line is printed)

3.4 Simple Wait

What can C print? Assume the child PID is 90210.

```
int main(void)
{
    pid_t pid = fork();
    int exit;
    if (pid != 0) {
        wait(&exit);
    }
    printf("Hello World\n: %d\n", pid);
}
```

Hello World 0
Hello World 90210 (or the failure case)

3.5 Exec

What will C print?

```
int main(void)
{
    char** argv = (char**) malloc(3*sizeof(char*));
    argv[0] = "/bin/ls";
    argv[1] = ".";
    argv[2] = NULL;
    for (int i = 0; i < 10; i++) {
        printf("%d\n", i);
        if (i == 3)
            execv("/bin/ls", argv);
    }
}
```

```
0
1
2
3
<output of ls>
```

3.6 Exec + Fork

How would I modify the above program using fork so it both prints the output of `ls` and all the numbers from 0 to 9 (order does not matter)? You may not remove lines from the original program; only add statements (and use fork!).

```
int main(void)
{
    char** argv = (char**) malloc(3*sizeof(char*));
    argv[0] = "/bin/ls";
    argv[1] = ".";
    argv[2] = NULL;
    for (int i = 0; i < 10; i++){
        printf("%d\n", i);
        if (i == 3) {
            pid_t pid = fork();
            if (pid == 0)
                execv("/bin/ls", argv);
        }
    }
}
```


4 Threads

4.1 Join

What does C print in the following code?

(Hint: There may be zero, one, or multiple answers.)

```
void *helper(void *arg) {
    printf("HELPER\n");
    return NULL;
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &helper, NULL);
    pthread_yield();
    printf("MAIN\n");
    return 0;
}
```

The output of this program could be "MAIN\nHELPER\n", "HELPER\nMAIN\n" or "MAIN\n". The actual order could be different each time the program is run. First, the `pthread_yield()` does not change the answer, because it provides no guarantee about what order the print statements execute in. Second, the helper thread may be preempted at any point (e.g., before or after running `printf()`). Last, the `main()` function can return without giving enough time for the helper thread to run, killing the process and all associated threads.

How can we modify the code above to always print out "HELPER" followed by "MAIN"?

Change `pthread_yield` to `pthread_join`.

```
void *helper(void *arg) {
    printf("HELPER\n");
    return NULL;
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &helper, NULL);
    pthread_join(thread, NULL);
    printf("MAIN\n");
    return 0;
}
```

4.2 Thread Stack Allocation

What does C print in the following code?

```
void *helper(void *arg) {
    int *num = (int*) arg;
    *num = 2;
    return NULL;
}

int main() {
    int i = 0;
    pthread_t thread;
    pthread_create(&thread, NULL, &helper, &i);
    pthread_join(thread, NULL);
    printf("i is %d\n", i);
    return 0;
}
```

The spawned thread shares the address space with the main thread and has a pointer to the same memory location, so i is set to 2. "i is 2"

4.3 Thread Heap Allocation

What does C print in the following code?

```
void *helper(void *arg) {
    char *message = (char *) arg;
    strcpy(message, "I am the child");
    return NULL;
}

int main() {
    char *message = malloc(100);
    strcpy(message, "I am the parent");
    pthread_t thread;
    pthread_create(&thread, NULL, &helper, message);
    pthread_join(thread, NULL);
    printf("%s\n", message);
    return 0;
}
```

"I am the child"

4.4 The Central Galactic Floopy Corporation

It's the year 3162. Floopies are the widely recognized galactic currency. Floopies are represented in digital form only, at the Central Galactic Floopy Corporation (CGFC).

You receive some inside intel from the CGFC that they have a GalaxyNet server running on some old OS called x86 Ubuntu 14.04 LTS. Anyone can send requests to it. Upon receiving a request, the server forks a POSIX thread to handle the request. In particular, you are told that sending a transfer request will create a thread that will run the following function immediately, for speedy service.

```
void transfer(account_t *donor, account_t *recipient, float amount) {
    assert (donor != recipient); // Thanks CS161

    if (donor->balance < amount) {
        printf("Insufficient funds.\n");
        return;
    }
    donor->balance -= amount;
    recipient->balance += amount;
}
```

Assume that there is some struct with a member **balance** that is **typedef**-ed as **account_t**. Describe how a malicious user might exploit some unintended behavior.

There are multiple race conditions here.

Suppose Alice and Bob have 5 floopies each. We send two quick requests: `transfer(&alice, &bob, 5)` and `transfer(&bob, &alice, 5)`. The first call decrements Alice's balance to 0, adds 5 to Bob's balance, but before storing 10 in Bob's balance, the next call comes in and executes to completion, decrementing Bob's balance to 0 and making Alice's balance 5. Finally we return to the first call, which just has to store 10 into Bob's balance. In the end, Alice has 5, but Bob now has 10. We have effectively duplicated 5 floopies.

Graphically:

Thread 1

```
temp1 = Alice's balance (== 5)
temp1 = temp1 - 5 (== 0)
Alice's balance = temp1 (== 0)
temp1 = Bob's balance (== 5)
temp1 = temp1 + 5 (== 10)
INTERRUPTED BY THREAD 2
```

RESUME THREAD 1

```
Bob's balance = temp1 (== 10)
THREAD 1 COMPLETE
```

Thread 2

```
temp2 = Bob's balance (== 5)
temp2 = temp2 - 5 (== 0)
Bob's balance = temp2 (== 0)
temp2 = Alice's balance (== 0)
temp2 = temp2 + 5 (== 5)
Alice's balance = temp2 (== 5)
THREAD 2 COMPLETE
```

It is also possible to achieve a negative balance. Suppose at the beginning of the function, the donor has enough money to participate in the transfer, so we pass the conditional check for sufficient funds. Immediately after that, the donor's balance is reduced below the required amount by some

other running thread. Then the transfer will go through, resulting in a negative balance for the donor.

Sending two identical `transfer(&alice, &bob, 2)` may also cause unintended behavior, since the increment/decrement operations are not atomic (though it is arguably harder to exploit for profit).

Since you're a good person who wouldn't steal floopies from a galactic corporation, what changes would you suggest to the CGFC to defend against this exploit?

The entire function must be made atomic. One could do this by disabling interrupts for that period of time (if there is a single processor), or by acquiring a lock beforehand and releasing the lock afterwards. Alternatively, you could have a lock for each account. In order to prevent deadlocks, you will have to acquire locks in some predetermined order, such as lowest account number first.

5 Interrupt Handlers

Refer to the “Pintos Interrupt Handler” section at the end of this discussion worksheet to answer these questions:

What do the instructions `pushal` and `popal` do?

They push and pop all the general-purpose 32-bit x86 registers onto/from the stack.

The interrupt service routine (ISR) must run with the kernel’s stack. Why is this the case? And which instruction is responsible for switching the stack pointer to the kernel stack?

The user program’s stack pointer may be invalid, or the user program could be using memory below the stack pointer. The CPU itself will switch the stack to the kernel stack (either because of an external interrupt, a trap, or a programmed interrupt). We do not need to write an instruction in the ISR to do this.

The `pushal` instruction pushes 8 values onto the stack (32 bytes). With this information, please draw the stack at the moment when “`call intr_handler`” is about to be executed.

ds
es
fs
gs
pushal’s 8 general purpose registers
pointer to (%esp + 4)

What is the purpose of the “`pushl %esp`” instruction that is right before “`call intr_handler`”?

It is a pointer to the part of the stack that contains all the registers. In pintos, this is accessed as the “`intr_frame`” struct.

Inside the `intr_exit` function, what would happen if we reversed the order of the 5 `pop` instructions?

The `pop` instructions need to be in their current order. They are exactly the reverse order of the corresponding push instructions, because our stack is First-In-Last-Out.

5.1 Pintos Interrupt Handler

```

1 /**
2  * An example of an entry point that would reside in the interrupt
3  * vector. This entry point is for interrupt number 0x30.
4  */
5 .func intr30_stub
6 intr30_stub:
7     pushl %ebp      /* Frame pointer */
8     pushl $0        /* Error code */
9     pushl $0x30     /* Interrupt vector number */
10    jmp intr_entry
11 .endfunc
12 /* Main interrupt entry point.
13
14    An internal or external interrupt starts in one of the
15    intrNN_stub routines, which push the 'struct intr_frame'
16    frame_pointer, error_code, and vec_no members on the stack,
17    then jump here.
18
19    We save the rest of the 'struct intr_frame' members to the
20    stack, set up some registers as needed by the kernel, and then
21    call intr_handler(), which actually handles the interrupt.
22
23    We "fall through" to intr_exit to return from the interrupt.
24 */
25 .func intr_entry
26 intr_entry:
27     /* Save caller's registers. */
28     pushl %ds
29     pushl %es
30     pushl %fs
31     pushl %gs
32     pushal
33
34     /* Set up kernel environment. */
35     cld                /* String instructions go upward. */
36     mov $SEL_KDSEG, %eax /* Initialize segment registers. */
37     mov %eax, %ds
38     mov %eax, %es
39     leal 56(%esp), %ebp /* Set up frame pointer. */
40
41     /* Call interrupt handler. */
42     pushl %esp
43 .globl intr_handler
44     call intr_handler
45     addl $4, %esp
46 .endfunc

```

```
48 /* Interrupt exit.
49
50 Restores the caller's registers, discards extra data on the
51 stack, and returns to the caller.
52
53 This is a separate function because it is called directly when
54 we launch a new user process (see start_process() in
55 userprog/process.c). */
56 .globl intr_exit
57 .func intr_exit
58 intr_exit:
59     /* Restore caller's registers. */
60     popal
61     popl %gs
62     popl %fs
63     popl %es
64     popl %ds
65
66     /* Discard 'struct intr_frame' vec_no, error_code,
67        frame_pointer members. */
68     addl $12, %esp
69
70     /* Return to caller. */
71     iret
72 .endfunc
```